T H E V A L U E A D D V C

How the Game of Venture Capital Really Works
and How to Win in the Age of AI

Trace Cohen

New York · 2026

Who This Book Is For

✔️ Founders deciding whether venture capital is right for their company
✔️ Emerging fund managers building their first or second vehicle
✔️ LPs and allocators seeking to understand what good fund management looks like
✔️ Startup employees trying to understand their equity and the game being played around them
✔️ Operators, advisors, and anyone building in the AI era who wants to understand the capital behind it
✔️ Anyone who has ever watched the venture world from the outside and wanted to understand how it actually works

How to Use the Interactive Tools

This book comes with a companion set of interactive calculators and tools at valueaddvc.com. Throughout the text, you'll see references to frameworks that can be modeled dynamically online. Look for the 🔧 icon for tool references.

First and foremost, this book is dedicated to my father, my dad, my business partner, and my best friend, Brian Cohen.

He has been the driving force behind my entrepreneurial journey for nearly two decades. Long before I understood what it truly meant to build something, he was already doing it—patiently and methodically.

We have built companies together. We have sold businesses together. We have invested together. We have traveled the world together.

And yes, I invited him to my bachelor party. That probably says enough.

Our partnership has never been just about business. It has been built on trust, loyalty, shared ambition, and a belief that building something meaningful matters. I would not be here without him.

This book is also dedicated to my mom, Carol.

My parents built their business together while raising three kids. We practically grew up in their office. Their employees helped take care of us. Work and family were never separate. They were part of the same life.

My mom has always been my confidant and has always had my best interests at heart. Much of my creativity comes from her. She showed me that you can be disciplined and imaginative at the same time. Structured and intuitive. Strong and caring.

Finally, to my wife, Sarah.

She is my rock. She supports my ambitions and puts up with my occasional ADHD side quests with patience and humor. She steadies me when I move too fast and reminds me what actually matters.

We met on Hinge, but we grew up a mile from each other. That still feels meaningful.

She is my partner in every sense of the word.

If this book reflects ambition, curiosity, and the desire to build something worthwhile, it is because of the three of them.

The Game

Venture capital is a game.

Not a metaphor. An actual game with rules, odds, players, and a scoreboard that most participants never fully see. The rules keep changing. The goalposts keep moving. Nothing is quite as it appears from the outside, and not much more than that from the inside.

Here’s what makes it different from most games: you don’t have to be a venture capitalist to be playing it. If you work at a startup, you’re playing it. If you’ve ever thought about starting a company, you’re about to play it. If you’re an investor, an allocator, an employee at a tech company, or simply someone trying to understand why the economy keeps getting reshaped by forces that nobody seems to explain clearly, you’re in the game whether you chose to be or not.

Most people play without knowing the rules.

That’s not an accident. The venture capital and startup world has its own language, its own logic, its own set of assumptions that insiders absorb over years and outsiders are rarely given directly. The result is a world that looks chaotic from the outside, where some founders raise millions on a napkin sketch and others with real businesses can’t get a meeting, when actually it operates by a clear, consistent, and learnable set of mechanics.

This book explains those mechanics. Not just for founders. Not just for investors. For everyone.

Because here’s the truth nobody says out loud: 99% of founders who raise venture capital probably shouldn’t. Not because they aren’t talented or hardworking or building something real, but because VC is a specific tool for a specific kind of company, and most companies aren’t that kind. Once you raise, you’re always raising. The expectations only go up. The game doesn’t pause.

But for the rare founder with a genuine unfair advantage, the right skills, the right network, the right obsession for the right problem at the right moment, venture capital is the greatest accelerant that exists. The question is whether you understand the game well enough to know if that’s you.

That’s what this book is for. Not to talk you into playing. Not to talk you out of it. To give you the map so that whatever you decide, you decide with your eyes open.

“

Venture capital is a game with constantly moving goalposts and expectations. Founders need to understand not only their own business but their investor’s business as well.

We invest in founders with unfair advantages. What can they do that no one else can? What do they know that others don’t? What makes them the one who figures this out?

Once you raise, you’re always raising.

Trace Cohen
Founder · Investor · Author

Preface

Why Everyone Needs to Understand This Game

I kept having the same conversation.

Not just with founders or investors, with everyone. Friends, family, colleagues, people at dinner parties who'd heard I was in venture capital and wanted to understand what that actually meant. Smart, curious people who felt like they were watching a game they'd never been taught to play. They saw the headlines, billion-dollar valuations, founders raising millions on a pitch deck, companies going from zero to IPO in five years, and they had the same question underneath all of it: how does this actually work?

Nobody had given them a straight answer. The venture world has always been better at mystifying itself than explaining itself. The jargon multiplies. The insiders speak in acronyms. The outsiders nod along and go home still confused.

This book is my attempt to fix that. For everyone.

My name is Trace. I was literally named after a supercomputer.

In 1988, one of my parents' clients was Multiflow Computer, a company founded by Josh Fisher and his team from Yale University. Their product was the Multiflow Trace—a minisupercomputer that pioneered a radical idea called Very Long Instruction Word (VLIW) architecture. The Trace could initiate up to twenty-eight operations in a single processor cycle by exploiting instruction-level parallelism: the ability to execute multiple computations simultaneously rather than one at a time. It was, in the language of computer science, parallel processing before parallel processing had a mainstream name.

My parents thought naming their son after a supercomputer was a perfectly reasonable thing to do. They were that kind of people.

Multiflow shipped roughly 120 systems before the company folded in the early 1990s. The product failed commercially. But the technology it invented—parallel processing, the idea that you could break complex problems into pieces and solve them all at once—powered virtually every high-performance computing system that followed. The GPUs that train modern AI models, the cloud infrastructure that runs every SaaS company in this book, the supercomputers that process the data behind every chart you'll see in these pages: all of them descend, in a direct technical lineage, from the principles the Multiflow Trace pioneered.

I was named after a machine that failed and whose idea won. If that isn't venture capital in a sentence, I don't know what is.

My parents were entrepreneurs. They built and sold a large tech PR agency in New York, running communications for all of Sony and all of IBM at a time when those names defined the center of gravity of the technology world. I

grew up on computers the way other kids grew up on baseball cards. I was in their commercials. I was their product tester. By the time I was in middle school, I could tell you the difference between a beta release and a launch build.

By the time I was seventeen, I was building companies.

My first started in my dorm room in 2007. My roommate and I built websites with Dreamweaver, HTML, and CSS, helping people get jobs online before SEO was a career and back when you needed a .edu email to create a Facebook account. That company went on a long, winding, educational journey before selling in 2021. Fourteen years. A lot of lessons.

Along the way I became member 400 of WeWork, back when the New York tech scene fit in a single building and everyone knew everyone. I joined New York Angels. I started investing in founders. I made good bets and bad ones, learned from both, and kept showing up.

Then AI changed everything. Not overnight. That is not how technology actually works, despite what headlines suggest. It changed gradually, then suddenly. I had been investing in AI companies for years, but over the last two years something shifted. I stopped just writing checks and started building with the tools myself, connected to Vercel, Supabase, Google Cloud, and a growing list of other services. Building under real constraints, with real users, making real mistakes, taught me more about AI-era company building than any pitch deck I've reviewed.

I've sat on both sides. As a founder I modeled the full dilution waterfall before signing any term sheet. As an investor I reserved capital for pro-rata in my best companies. The ones where I exercised at Series A are where the math worked out.

But I didn't write it for insiders. I wrote it because the game affects everyone: founders deciding whether to raise, investors learning how to think, employees trying to understand their equity, parents watching their kids enter this world, and anyone trying to make sense of the forces reshaping the economy around them.

You don't have to be in venture capital for this book to matter. You just have to live in a world that venture capital is changing, which, at this point, is everyone.

Before we get into the frameworks, a note on language. Venture capital has its own vocabulary, its own shorthand, its own set of terms that insiders use fluently and outsiders are rarely taught directly. I've included a comprehensive Venture Capital Glossary and Metrics Reference in Appendix A at the back of this book. I define every term in these pages there, in plain English. If you encounter a term you don't recognize, flip to the back. No one should have to pretend they know what "DPI" means to follow the argument. That said, you'll pick up most of the vocabulary naturally as you read—I introduce each term in context, the first time it appears, with enough explanation that you won't need to flip anywhere unless you want to.

How to Read This Book

I've structured this book in five parts, each building on the last. Part I explains the macro environment—what happened when money got expensive. Part II maps where value actually accrues in AI. Part III covers fund mechanics and LP dynamics. Part IV is the operator's manual for anyone in the middle of building or raising right now. Part V is about what endures across cycles.

You can read straight through or jump to whatever is most relevant to where you are right now. Each chapter stands on its own, though the cumulative effect of reading in order is that you'll develop the kind of structural intuition that takes most practitioners years to build.

A few things to know before you start:

The gray-boxed stories throughout are real. I've changed names and details where confidentiality requires it, but the mechanics and lessons are exact. The pull quotes at the end of each chapter are the one-line versions of the chapter's core argument—useful for quick reference or for the kind of person who reads the last page of a mystery first (no judgment, I do it too).

The figures and charts make the math visual. If you're a numbers person, they'll sharpen your understanding. If you're not, they'll make the concepts accessible without requiring you to love spreadsheets. (Though after reading Chapter 11, you might start.)

Throughout the book, you'll find Reader Challenges—questions designed to test your assumptions before the data arrives. I encourage you to actually answer them. The gap between what you assume and what the data shows is where the deepest learning happens. Scribble in the margins. Argue with me. That's what margins are for.

🎲 The VC Bingo Card

Things you will absolutely hear in LP meetings, pitch sessions, and board calls. Check them off as you go. Five in a row wins nothing except the knowledge that venture capital is predictable in all the ways it claims not to be.

"We're platform agnostic"	"Our DPI speaks for itself"	"We add value beyond capital"	"It's a $100B TAM"	"We're seeing strong signals"
Someone mentions "first principles"	"We're in the early innings"	"Our fund is right-sized"	"The founder has a unique insight"	"We passed on [famous company]"
"We have proprietary deal flow"	"We invest in category-defining companies"	FREE SPACE: "AI-native"	"We focus on founder-market fit"	"The market is at an inflection point"
"Our portfolio has strong markups"	PowerPoint has 80+ slides	"We're contrarian"	Someone references Sequoia	"We have a unique sourcing advantage"
"This is a winner-take-all market"	"It's like [X] for [Y]"	"Our portfolio NPS is very high"	"We move fast on conviction"	"The timing is perfect"

Pro tip: If you check off the entire card in a single meeting, you are either at a very productive conference or a very unproductive LP meeting. Either way, you should leave.

Quick Reference: Key Terms

The vocabulary you'll encounter most often in these pages. The complete Venture Capital Glossary and Metrics Reference is in Appendix A.

LP	Limited Partner. The institution or individual who invests into a venture fund. Endowments, foundations, pension funds, family offices.
GP	General Partner. The fund manager who makes investment decisions and deploys LP capital.
VC	Venture Capital. Private equity focused on early-stage, high-growth companies in exchange for equity.
MOIC	Multiple on Invested Capital. How many times you got back what you put in. 3x MOIC = $3 returned per $1 invested. The most intuitive performance measure.
DPI	Distributed to Paid-In. Cash actually returned to LPs as a multiple of invested capital. Unlike TVPI, DPI only counts money in the bank.
TVPI	Total Value to Paid-In. DPI plus current estimated value of unrealized investments. A useful snapshot, but subject to mark assumptions.
IRR	Internal Rate of Return. Time-weighted return accounting for when cash flows in and out. Early returns are worth more than equivalent late returns.

NRR	Net Revenue Retention. Revenue kept from existing customers year-over-year including expansion. Above 100% means growth without adding new customers.
ARR	Annual Recurring Revenue. Annualized value of recurring subscription contracts. The heartbeat metric for SaaS businesses.
Burn Multiple	Net cash burned ÷ net new ARR. How efficiently are you buying growth? Under 1x is elite. Over 3x is a problem.
D/D	Due Diligence. The research process investors use before making an investment. Financial, legal, technical, reference-based.
Cap Table	Capitalization Table. The full breakdown of who owns what: founders, employees (via options), and all investors.
Carry	Carried Interest. The fund manager's share of profits, typically 20% of gains above the hurdle rate. How GPs earn from strong performance.
Pro-Rata Rights	The right to invest in future rounds to maintain your ownership percentage. One of the most valuable rights an early investor holds.
SAFE	Simple Agreement for Future Equity. Converts to equity at a future priced round without the overhead of debt structure. The most common early-stage instrument.
Down Round	A financing round at a valuation below the prior round. Once rare enough to be newsworthy; now a structural reality in repriced markets.
Dry Powder	Capital committed to a fund but not yet deployed. In contractions, the managers with dry powder write the best deals.
J-Curve	The pattern of negative early returns (from fees and write-offs) followed by positive returns as exits materialize. Every fund goes through it.
Lead Investor	The investor who sets the terms, writes the largest check, and typically joins the board. Everyone else follows their lead.
Liquidation Preference	The right to receive a specified return before common shareholders in an exit. 1x non-participating is standard and fair. Anything more, read twice.
Vintage Year	The year a fund makes its first investment. Used for peer comparison because market conditions affect every fund in a cohort.
Rule of 40	Revenue growth rate + EBITDA margin. A sum above 40% is

	considered healthy at growth stage. Simple, useful, and hard to fake.
Runway	The number of months a startup can operate before running out of cash at current burn. Under 12 months is a warning sign. Under 6 is a crisis.
Term Sheet	A non-binding document outlining the key terms of an investment. The first draft sets the negotiation anchor. Read every line.
Markup	An increase in valuation from one round to the next. Paper markups make LPs feel good. DPI makes them come back.
Bridge Round	A smaller financing meant to extend runway between priced rounds. Sometimes strategic. Sometimes a sign that the next real round isn't coming.
Follow-On	Additional capital invested in an existing portfolio company in subsequent rounds. The best fund strategies protect ownership through disciplined follow-on.
Secondaries	The buying and selling of existing fund interests or company shares on a secondary market. Increasingly important for LP liquidity and GP flexibility.
GP Commit	The amount a fund's general partners invest in their own fund. Typically 1–3% of fund size. Signals alignment. LPs notice when it's low.
Power Law	The distribution pattern where a small number of investments generate the vast majority of returns. In venture, the top 5% of deals drive most outcomes.

Contents

The Game
Preface: Why Everyone Needs to Understand This Game
Key Terms at a Glance
Introduction: The Repricing

PART I — The Repricing of Venture
Chapter 1: Capital Cost, Liquidity, and Duration
Chapter 2: Building With AI: What I Learned by Actually Doing It
Chapter 3: The Bifurcation: Mega-Funds and Emerging Managers
Chapter 4: Exit Distribution: Frequency vs. Visibility
Chapter 5: LP Capital Allocation Under Constraint

PART II — Where Value Accrues in the AI Era
Chapter 6: The AI Stack: Where Smart Money Actually Goes
Chapter 7: Why Vertical AI Wins
Chapter 8: The Vertical AI Moat Stack
Chapter 9: Vertical AI Economics
Chapter 10: Building for M&A vs. IPO

PART III — The Capital Stack
Chapter 11: Emerging Managers: Ownership, Asymmetry, and Edge
Chapter 12: How LPs Actually Think and How to Raise From Them
Chapter 13: LP Due Diligence: What Allocators Actually Evaluate
Chapter 14: Fund Size Drift and Return Compression

PART IV — The Operator’s Manual
Chapter 15: Fundraising Is a Funnel
Chapter 16: The Series A Gap Is Real
Chapter 17: Burn Discipline in a Repriced Market
Chapter 18: Capital Stack Sequencing
Chapter 19: Founder-Investor Alignment

PART V — Structural Advantage
Chapter 20: Capital Efficiency in the AI Era
Chapter 21: Reading the Market Cycle
Chapter 22: The Discipline of Structure

Closing Note
Afterword: The Data Has Arrived
Appendix A: Venture Capital Glossary and Metrics Reference
Appendix B: Scenario Models
Appendix C: Suggested Reading
Index
About the Author

Introduction

In 1997, my parents were in the room.

They ran a technology PR agency in New York, one of the best in the country, and IBM was their client. The campaign they helped build was for a chess match between Garry Kasparov, the greatest human player alive, and Deep Blue, IBM's supercomputer. Human versus machine. The first contest of its kind that the world took seriously.

Kasparov lost. The world was briefly unsettled. My parents came home and told me about it the way other parents might have described watching a moon landing.

I was named after a computer. Not Deep Blue, but another machine entirely: the Multiflow Trace, a VLIW minisupercomputer that pioneered parallel processing—the ability to execute dozens of operations simultaneously. My parents ran PR for the company that built it. The Trace shipped about 120 units before Multiflow folded. The company died. The technology it invented now powers modern computing—from the GPUs training today's AI models to the cloud infrastructure running every startup in this book. I grew up testing products, appearing in commercials, learning the difference between a beta release and a launch build before most of my classmates knew what a hard drive was. By the time I was in middle school I could tell you why IBM's stock moved on a press release. By seventeen I was building websites in a dorm room and charging for them.

I'm telling you this because it explains something about why I wrote this book and who I think it's actually for.

The venture capital industry has always had a Deep Blue problem. The insiders understand the machine. The outsiders watch the match and feel vaguely unsettled by something they can't quite name. The machine makes moves that look either brilliant or arbitrary depending on whether you know how it thinks. And almost nobody explains how it thinks.

That changes here.

The Repricing

In early 2022, venture capital did not collapse. It recalibrated.

For more than a decade, capital had been inexpensive and abundant. Public markets rewarded growth over everything else. Private markets extended that logic further, and further still. Fund sizes expanded. Rounds moved faster. Valuations floated upward on anticipation of future scale rather than present reality. Founders were told to grow at all costs. Investors were told the window would stay open.

Then the cost of capital changed.

Interest rates rose at the fastest pace in decades. Public technology multiples compressed. The IPO market slowed, then effectively froze. Liquidity that once felt cyclical began to feel conditional. Institutional allocators watched their public portfolios decline while private marks hadn't caught up yet. The denominator effect kicked in. Commitment pacing tightened across the board.

Nothing about entrepreneurship disappeared. Builders kept building. Markets kept evolving. But the environment that had supported aggressive long-duration bets changed abruptly, and venture, an asset class that depends entirely on future exits it cannot control, felt it immediately.

This was not a moral correction. It was a financial one.

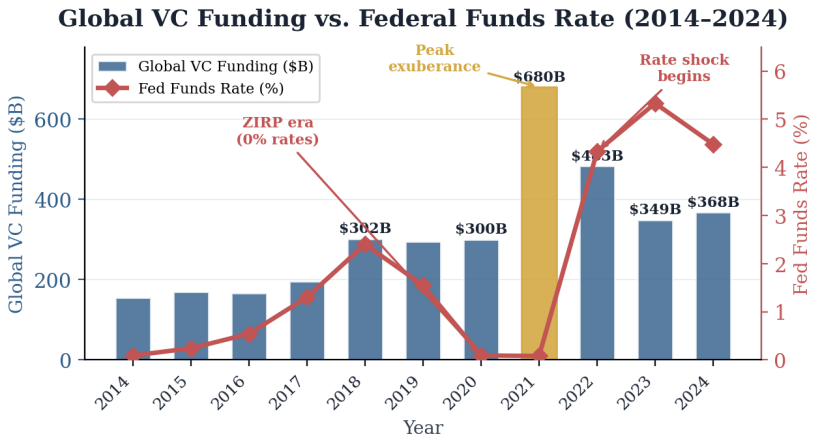

When money was free, capital flooded in. When rates rose, the denominator effect kicked in. Sources: KPMG, Crunchbase, Federal Reserve

Global VC funding vs. federal interest rates (2014–2024). The relationship between rate policy and venture deployment is not coincidence; it's arithmetic.

The repricing clarified something that had been invisible for a decade under cheap money: the rules of venture capital had not changed. They had simply been buried. When the tide went out, the mechanics were exposed.

> ***Cycles rotate. Structure compounds. The next decade will reward those who understand that difference.***

Whether you're raising your first fund, building your first company, trying to understand your equity as an early employee, or simply watching from the outside and wanting to understand what actually drives these decisions, the mechanics are learnable. You don't need to be an insider. You need to understand how the machine thinks.

Let's get into it.

P A R T I

The Repricing of Venture

Capital, rates, exits, and the truths that got exposed when the money got expensive

Key Themes

Why the cost of capital changes everything about venture math

What building with AI taught me that investing never could

How fund size determines probability of success

Where exits actually happen vs. where headlines say they do

Why your LP said no even when they loved your fund

What You'll Learn

How to calculate the present-value impact of rate changes on private company valuations

The specific diligence questions that separate informed AI investors from tourists

Why a $75M fund and a $750M fund targeting the same exits produce radically different results

The denominator effect: the mechanical reason your best LP meeting still ended in a no

How to read the exit distribution data that reveals the real odds of venture

---

📝 Reader Challenge: Before reading Part I, write down what you think the average time-to-exit is for a venture-backed company. Then write down what percentage of exits happen via IPO vs. acquisition. Check your answers against the data in Chapters 1 and 4. The gap between your guess and reality is exactly why this section exists.

Chapter 1

Capital Cost, Liquidity, and Duration

The day the math changed

Venture capital is a duration asset. That's the most important frame for understanding everything that happened in 2022 and after.

When you make a venture investment, you're not getting a coupon every six months. You're making a bet that some company, at some point in the future, maybe in seven years, maybe in eleven, will be worth enough that your ownership stake returns multiples of what you put in. The present value of that future outcome depends entirely on the cost of capital today.

When rates are near zero, future cash flows discount gently. When rates rise fast, as they did starting in 2022, that same future company is worth materially less in present-value terms. It's a basic formula that ran on every terminal on Wall Street simultaneously. The market didn't wait for anyone to catch up.

I watched this happen in real time across my portfolio. Companies that had been marked at strong multiples in 2021 started showing cracks in their next fundraises, not because the businesses had changed, but because the discount rate had. The math had moved beneath them while they were busy executing.

> ***Capital is never free. It's priced differently across cycles. Know which cycle you're in before you make your next bet.***

> 💬 **TRACE'S SASSY SIDEBAR #1**
>
> *Translation: if someone tells you "money is cheap right now," ask them what year it is. Because that sentence has an expiration date, and it's usually about eighteen months before you hear it.*

What the Repricing Actually Revealed

Late-stage valuations corrected sharply. Companies that raised at 30x revenue found their next rounds pricing at 8–10x. Down rounds came back as a structural necessity. Liquidity windows narrowed dramatically, with IPO issuance falling more than 80% from its 2021 peak. Institutional LPs adjusted their commitment pacing through the denominator effect: when public

portfolios decline, private allocations mechanically rise as a percentage, forcing LPs to slow new commitments even when they believe in the managers asking.

7–11 yrs
Average time to exit, post-repricing

80%+
Decline in US IPO volume from 2021 peak

None of this was permanent. Cycles turn. But the managers and founders who treated these mechanics as temporary aberrations made systematically worse decisions throughout the period. Understanding duration risk changes how you price rounds, plan reserves, and counsel founders on cash management.

The repricing was not a punishment. It was a recalibration to the historical mean. The managers who recognized this early preserved capital, maintained ownership discipline, and positioned themselves for the vintage-year opportunity that contractions always create. Those who waited for the old environment to return are still waiting.

The chapters that follow break down how each of these forces—cost of capital, exit distribution, fund mechanics, LP behavior—interact in practice. None of them operate in isolation. Together, they define the structural reality that every founder, investor, and allocator is navigating right now.

Chapter 2

Building With AI: What I Learned by Actually Doing It

What fifteen years of investing couldn't teach me, and six months in the terminal did

I want to tell you something I've hesitated to say in more formal contexts: I didn't truly understand AI until I started building with it myself.

I had been investing in AI companies for years. I sat on boards, reviewed metrics, asked the right questions in diligence. But writing the code—or at least prompting the AI that writes the code—teaches you something no deck or quarterly update ever will.

Learning in the Terminal

I started using Claude in the terminal. Not through a web interface, but connected directly through APIs, hooked up to Vercel for deployment, Supabase for the database layer, Google Cloud for infrastructure, and a growing ecosystem of other services. I was building things that broke constantly and required me to understand why.

The learning curve was steep in the best possible way. When you're debugging an agent loop running in a terminal, connected to a live database, you develop an intuitive understanding of what these systems can and cannot do that no amount of reading provides. You learn where they hallucinate. You learn where they're remarkably reliable. You learn what good prompting actually means in practice, not as a parlor trick, but as a real engineering discipline.

> *There's a moment every investor in this space needs to experience: asking an AI system to do something that requires real reasoning, watching it get partway there, debugging why it failed, and iterating until it works. That process teaches you more about AI capabilities and limitations than a year of attending conferences.*

Building ValueAddVC.com

The project that crystallized everything was building ValueAddVC.com, a platform giving founders, investors, and aspiring VCs the tools and frameworks usually locked behind years of experience or expensive relationships.

I built it module by module. A deal memo generator. A portfolio tracking tool. An LP communication framework. A fund economics calculator. Each module taught me something different: where AI could operate without supervision, where human judgment was irreplaceable, what working software actually meant when AI was writing 80% of the code.

I made mistakes constantly. I shipped things that didn't work. I iterated in public, which was uncomfortable but accelerating. The feedback loop of building something real, that actual users could break in ways I hadn't anticipated, compresses learning in a way nothing else does.

What I Got Wrong

I got the timeline wrong. I thought certain capabilities were further away than they were, particularly the ability of AI to maintain context over long tasks and handle genuinely complex reasoning chains. I was repeatedly surprised by what was possible with the right prompting and architecture.

I got the failure modes wrong too. I expected AI to fail loudly and obviously. The actual failure patterns are quieter and more dangerous: confident wrong answers, plausible-sounding hallucinations, gaps in reasoning that look like completions. Learning to spot these, to know when a model is extrapolating versus reasoning, took time I couldn't have shortcut.

And I underestimated the importance of the human layer. AI is extraordinarily capable at generating output. It is much less capable at knowing what output is worth generating. That judgment, editorial, strategic, aesthetic, remains deeply human. The best AI-assisted work I've seen combines AI throughput with human taste in ways neither could achieve independently. The human is not the bottleneck. The human is the point.

What It Changed About How I Invest

Building with these tools changed my investment criteria in concrete ways. I evaluate vertical AI companies with much higher confidence because I've felt the integration friction myself. I understand what good AI product architecture looks like from the inside. I ask better diligence questions because I know which failures are common and which are red flags.

The most valuable investors in the AI era will not be those who understand AI in the abstract. They will be those who understand it from experience, who have felt the difference between a model that reasons well and one that merely sounds like it does.

> 💬 **TRACE'S SASSY SIDEBAR #2**
>
> *I once spent three hours debugging a Supabase auth loop at 2am on a Tuesday. That single session taught me more about AI product architecture*

> *than six months of LP update calls. If your AI thesis starts with "we believe" and not "we built," we're going to have a conversation.*

If you invest in AI companies and haven't spent significant time building with AI tools, start. Not to become an engineer. To become a better investor.

🔍 AI DILIGENCE QUESTIONS TO ASK IN YOUR NEXT MEETING

1. What happens to your product if the underlying foundation model improves by 10x?
2. What proprietary data do you generate through usage that competitors cannot replicate?
3. How would a customer replace you? What would break in their workflow?
4. What percentage of your value comes from the AI layer vs. the workflow layer?
5. Show me the usage data. Not the signups—the repeat sessions and time-in-app.

These five questions separate investors who talk about AI from investors who understand what they're underwriting. Use them. The founders worth backing will have answers that surprise you.

> ***You cannot invest well in what you don't understand from the inside. Build something. Break something. The education is worth more than any conference.***

The AI era rewards a new kind of investor: one who can evaluate not just a pitch deck but a training pipeline, not just a TAM slide but an inference cost structure, not just a team's resume but their rate of shipping. The frameworks that follow will give you the tools to do exactly that. But the foundation—the willingness to build, to break, to learn from the inside—starts here.

In the next chapter, we'll look at what happens when the math of fund size meets the reality of exit distribution. Spoiler: most funds are too big for the outcomes the market actually produces.

The bifurcation is accelerating. In January 2026, Andreessen Horowitz alone raised $15 billion across new funds. That same period, all first-time venture funds in the U.S. combined raised $6.6 billion. Five AI companies captured $192 billion in cumulative venture funding—one-third of all AI investment globally. Deals above $500 million now account for 47% of all U.S. venture activity, up from 30% in 2010. The capital is concentrating at both ends: massive platforms and disciplined emerging managers. The middle continues to hollow out.

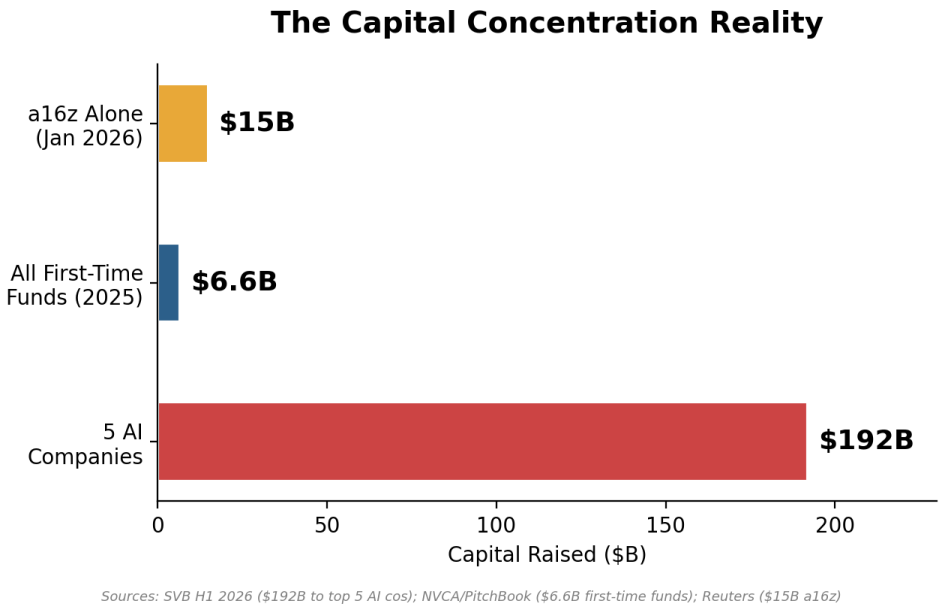

The capital concentration reality. A single firm raised more than double what all first-time funds combined raised in 2025. Sources: SVB H1 2026, NVCA/PitchBook, Reuters.

This is the structural reality emerging managers must navigate. You are not competing for the same capital as Andreessen Horowitz. You are not competing for the same deals as SoftBank. But the LP dollars that flow to those platforms are dollars that do not flow to you. The bifurcation creates two distinct games with two distinct sets of rules. The mistake is playing one while thinking you're in the other. The managers who thrive in this environment are those who understand exactly which game they're playing—and size their funds, their ownership targets, and their return expectations accordingly.

Chapter 3

The Bifurcation: Mega-Funds and Emerging Managers

Why the middle of the market got squeezed, and why that's actually good news for smaller funds

As venture capital scaled through the 2010s, the industry bifurcated. At the top, a small number of large platforms accumulated enormous capital. At the bottom, a wave of first and second-time managers launched focused, smaller vehicles. The middle, funds between $200M and $600M, got compressed.

This isn't random. It's a direct consequence of the math.

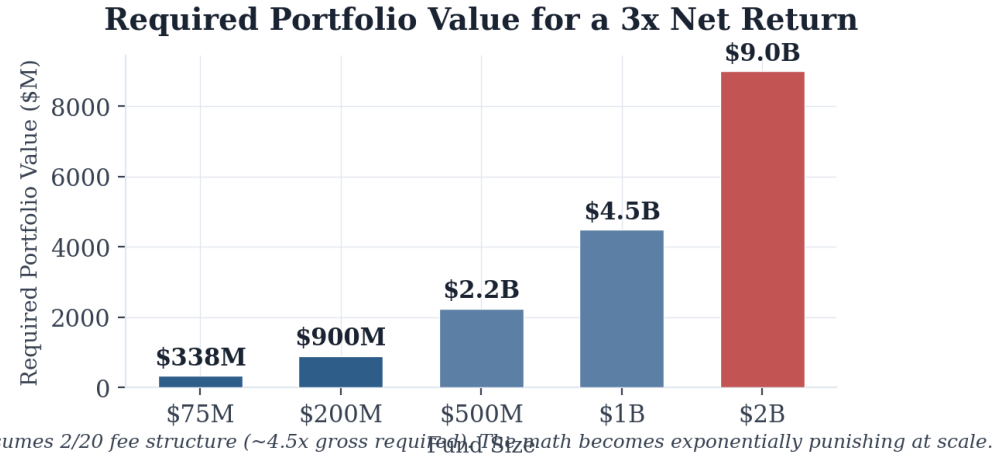

Required portfolio value for a 3x net return, by fund size. The math becomes punishing at scale, and the exit market doesn't expand proportionally with fund size.

A $2 billion fund targeting 3x gross needs to generate $6 billion in portfolio value. At 8–12% ownership at exit, that requires multiple companies reaching $3–5 billion-plus exits. In any given year, the global cohort of venture-backed exits above $5 billion might be fewer than twenty companies. You need to own several of them.

A $75 million fund has a fundamentally different problem. It needs roughly $225 million in distributions. With disciplined ownership of 8–12% at exit, a single $500 million acquisition can contribute $40–60 million. Two outcomes like that and you're most of the way home. $500 million acquisitions happen regularly.

I've had this conversation with first-time managers tempted to raise bigger than they should because 'bigger sounds more legitimate.' I tell them the same thing: your job isn't to raise the most capital possible. It's to raise the

amount your deal access, ownership levels, and exit distribution can actually support. Those are very different numbers.

> ***Fund size is not prestige. It is probability.***

> 💬 **TRACE'S SASSY SIDEBAR #3**
>
> *Every time a first-time manager tells me they're "thinking about $200M for Fund I because LPs keep offering," I ask them to name five $500M+ exits in their deal pipeline. The silence that follows is the whole lesson.*

The bifurcation also reshapes how LPs think about portfolio construction. A decade ago, an institutional allocator might have committed to fifteen funds across a range of sizes and felt adequately diversified. Today, the barbell is more deliberate: large platforms for market coverage and brand-name co-investment access, paired with small focused vehicles for alpha generation. The middle gets squeezed from both sides.

For emerging managers, this is the strategic reality: you are not competing with mega-funds for the same deals. You are competing for the same LP dollars. The pitch is not "we are smaller and therefore better." The pitch is "our size is the strategy." Every check we write can move the needle on fund returns. Every company we back can reach the ownership levels where outcomes actually compound. We are not a scaled-down version of Andreessen. We are playing a fundamentally different game with fundamentally better math.

That distinction matters because the data supports it. Funds under $250 million have historically generated higher median net IRRs and TVPIs than their larger counterparts across nearly every vintage year since 2005. The reasons are structural, not accidental: earlier entry points, higher ownership at lower valuations, and a broader exit universe that does not require billion-dollar outcomes to generate strong returns. When you need a $300 million exit to return 4x on a position, the math works in your favor. When you need a $3 billion exit, you are dependent on a handful of outcomes that may not materialize in your fund's lifecycle.

Chapter 4

Exit Distribution: Frequency vs. Visibility

What the headlines get wrong about how venture actually returns money

Ask most people what a successful venture exit looks like and they'll describe an IPO. A ringing bell. A founder in a hoodie. That's not wrong. It's just rare—and it's responsible for a significant amount of bad planning across the ecosystem.

The majority of venture-backed exits, across every market environment, are acquisitions. Strategic buyers pay for distribution, talent, and competitive defense. Many exits happen between $100M and $1 billion. They don't get TechCrunch coverage. They get a press release and a quiet congratulatory Slack from corp dev.

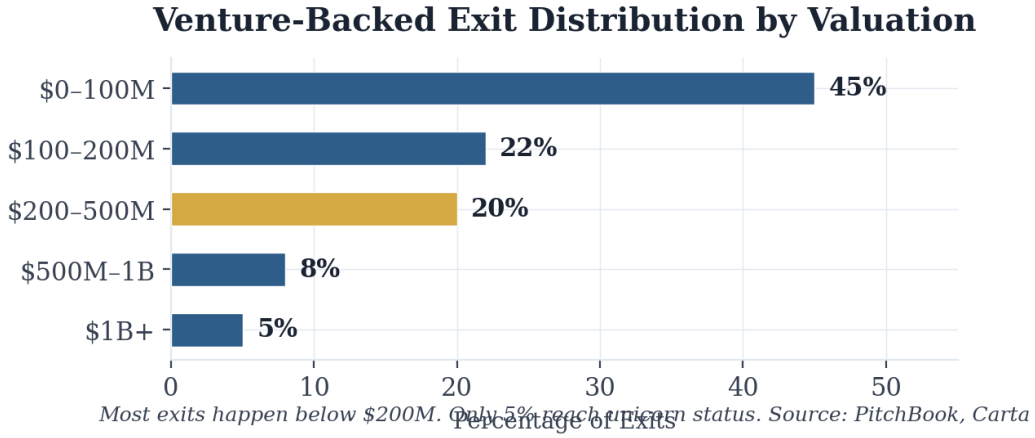

Exit distribution by valuation range. The $200M–$500M band represents nearly one in four exits. You don't read about most of them.

> ***Distribution curves shape venture more than headlines do. The exits that move your fund are rarely the ones on the front page.***

The Narrative Gap

At the beginning of a fund, most managers project three of twenty-five companies will achieve an outsized exit. LPs nod. The empirical reality: zero to one reach $1B+; four to six reach $200–500M; thirteen to fifteen exit below $100M or fail. The gap between projection and reality is the narrative gap. The managers who outperform close it by modeling conservatively—building funds that work under realistic distributions.

Chapter 5

LP Capital Allocation Under Constraint

Why your investors can't always say yes, even when they want to

One of the hardest things to explain to first-time fund managers is why a strong LP meeting sometimes still ends in a 'no.' The fund is performing. The strategy is clear. The manager is excellent. And yet.

The answer is usually not about the manager at all. It's about the LP's allocation mechanics: portfolio-level constraints that govern whether any new commitment is even possible, regardless of quality.

The Denominator Effect

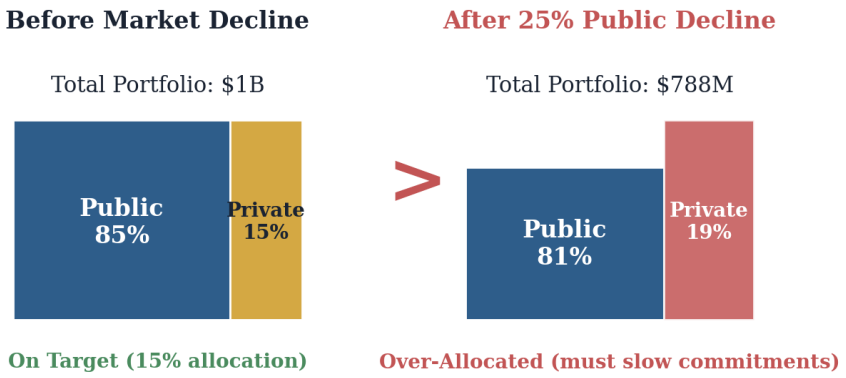

Private holdings didn't change. The denominator did. LPs must slow commitments regardless of manager quality.

The denominator effect. When public markets fall, private allocations rise as a percentage of total portfolio, mechanically, automatically, forcing LPs to slow commitments even when they love the manager.

Here's how it works: an endowment targets 15% allocation to private markets. When public equities drop 25%, the denominator shrinks. The private portfolio, still marked at prior values, now represents 19% of the total. They're suddenly over-allocated, and new commitments push them further over. So they pause. Regardless of whether your fund is the best thing they've seen all year.

I've been on the receiving end of this. You get a 'no' with a very warm tone and a promise to revisit next vintage. Once you understand LP mechanics, you stop taking it personally and start planning around it: raising before the freeze, not during it.

> ***Understanding your LP's constraints is as important as understanding your own portfolio. Capital responds to mechanics, not just to merit.***

❓ POP QUIZ — Chapter 5

1. What is the "denominator effect" and why does it cause LPs to pause commitments even when they love a manager?
2. True or false: The 2022 repricing was caused by startups building worse products.
3. A $2B fund targeting 3x needs how much in portfolio value? How many $5B+ exits happen globally per year?
4. What percentage of venture-backed exits happen via IPO vs. acquisition?
5. Why do most institutional LPs have a target allocation of only 5–15% to venture capital, even when historical returns justify more?

Answers (no peeking): 1. When public markets fall, private allocations rise as a % of total portfolio, mechanically pushing LPs over their target allocation. 2. False. It was a financial recalibration driven by rising interest rates and compressed public multiples. 3. $6B in portfolio value. Fewer than 20 exits above $5B per year globally. 4. The vast majority are acquisitions. IPOs make the headlines; M&A makes the distributions. 5. Illiquidity risk, J-curve drag, and the need for diversification across asset classes. Venture is high-return but also high-variance and long-duration.

Part I at a Glance

Chapter 1	Capital is never free. Rising rates repriced the entire venture landscape.
Chapter 2	AI changed the tools, the speed, and the cost structure. Investors who build understand what they fund.
Chapter 3	Fund size is probability, not prestige. The math favors focused, right-sized funds.
Chapter 4	Most exits happen through acquisition, not IPO. Exit distribution data should drive portfolio strategy.
Chapter 5	LP behavior is driven by allocation mechanics, not just conviction. The denominator effect is structural.

Key Numbers from Part I

$621B	<20	8–12%	$6B
Peak global VC funding (2021)	Exits above $5B per year globally	Target ownership at exit for disciplined funds	Portfolio value a $2B fund needs for 3x

🔧 *Model these numbers yourself at valueaddvc.com/calculator*

📝 Notes & Scribbles

After reading Part I, what's the single biggest assumption about venture capital you held before this book that the data challenged? Write it down. Then write what you believe now.

P A R T I I

Where Value Accrues in the AI Era

Not all AI is created equal. Here's where the real money gets made.

Key Themes

Why the most obvious AI investments are often the worst ones

The compounding logic of vertical over horizontal

How switching cost creates durable value

Reading the economics that separate winners from hype

Building for the exit that will actually happen

What You'll Learn

How to evaluate which layer of the AI stack is actually defensible vs. commoditizing

The five-layer moat stack that separates durable vertical AI from hype

Why 120% NRR changes the entire risk profile of an investment

How to read burn multiples and know when efficiency signals real discipline vs. starvation

The strategic difference between building for M&A vs. IPO—and why optionality matters more than ambition

---

📝 Reader Challenge: Pick an AI company you know—one you've built, invested in, or simply admire. Before reading Part II, write down which layer of the AI stack it operates at and how many moat layers it has. Be honest. Revisit after Chapter 8.

Chapter 6

The AI Stack: Where Smart Money Actually Goes

A map of the ecosystem and why the most obvious investments are often the worst ones

When AI became the dominant venture theme, money flooded in at every layer of the stack. Foundation model companies raised billions. Infrastructure tooling raised hundreds of millions. Orchestration frameworks, vector databases, fine-tuning platforms: everything touched by AI attracted capital at extraordinary speed.

Most of that capital will not generate the returns investors expected. Not because AI isn't transformative, it clearly is, but because investors confused 'this technology is important' with 'this company is defensible.' Those are very different claims.

The AI Stack

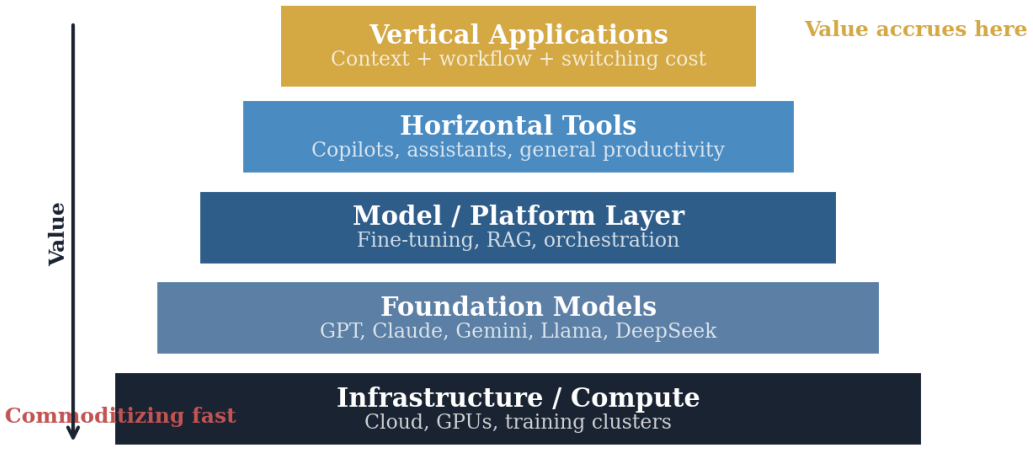

Infrastructure commoditizes. Context compounds. Invest accordingly.

The AI stack. Value migrates upward over time, from infrastructure (commoditizing fast) to the application layer (where context, workflow, and switching cost accumulate).

Foundation Models: The Infrastructure Trap

Foundation model companies are remarkable technical achievements. But as a venture investment thesis they carry a structural problem: training costs require sustained capital most venture funds can't support at scale. Competitive dynamics are intense among well-capitalized players. Pricing

pressure is relentless as compute costs fall and open-source alternatives improve. The companies building on top of the models often capture more durable value than the model providers themselves.

The Application Layer: Where Emerging Managers Win

For most emerging managers, the application layer is where the risk-return profile makes the most sense. Application companies that embed deeply into specific vertical workflows accumulate domain context, proprietary data, and switching cost that create genuinely durable businesses. The foundation model underneath them may change. The orchestration layer may be replaced. But the workflow that's been rebuilt around their product, the compliance infrastructure in their deployment, and the years of domain-specific training data they've accumulated cannot easily be replicated by a competitor with access to the same underlying model.

AI Categories: Defensibility vs. Competitive Intensity

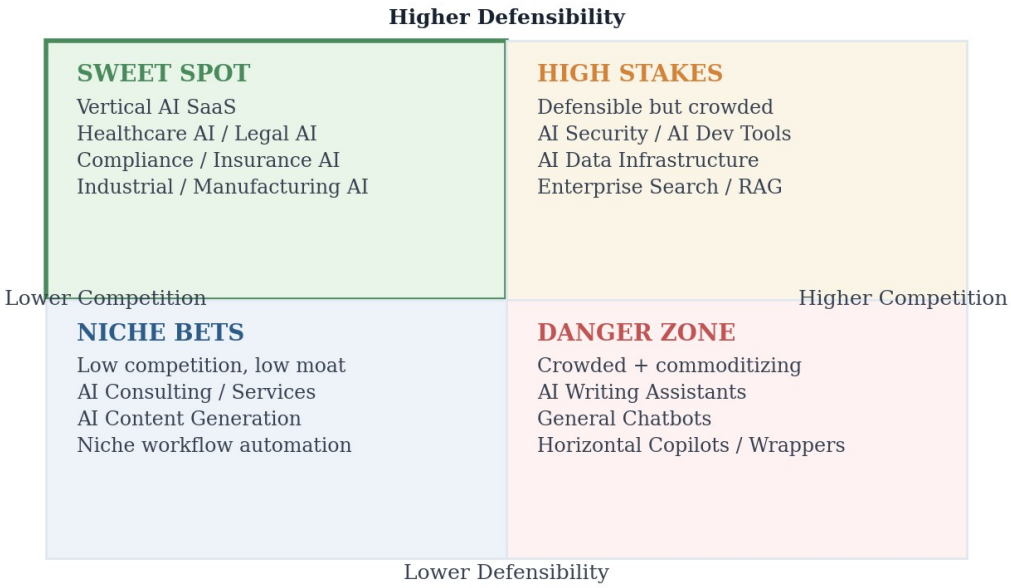

Emerging managers should concentrate capital in the upper-left quadrant.

AI categories by defensibility vs. competitive intensity. The upper-left quadrant, high defensibility, lower competitive density, is where emerging managers should concentrate capital.

This is where the AI investment thesis converges with the emerging manager math. A vertical AI company with embedded workflows, proprietary data, and high switching cost doesn't need a $5 billion exit to move a $75 million fund. It needs to be genuinely hard to displace, and that kind of durability starts at the application layer, not the infrastructure layer.

> ***Models commoditize. Context compounds. Invest accordingly.***

Chapter 7

Why Vertical AI Wins

The compounding logic of domain specificity, and why breadth is a trap

Horizontal AI tools promise to do everything for everyone. Vertical AI tools promise to do one thing extraordinarily well for a specific industry. Chapter 6 mapped the landscape. This chapter makes the investment case, because understanding why vertical wins changes not just what you invest in, but how you evaluate every AI company you see.

Foundation model access is becoming available to everyone. When everyone can access equivalent model capability, differentiation migrates away from the model and toward everything around it: domain knowledge, workflow integration, accumulated data, compliance infrastructure. All of those things take years to build. None can be replicated by a competitor who just got API access to the same model you're using.

Vertical AI companies compound in ways that horizontal tools simply cannot. Their proprietary data becomes more valuable over time: every processed claim, dispatched truck, or underwritten policy trains the system in ways a new entrant can't replicate without years of equivalent volume. Their workflow integration deepens. When the product sits inside the EHR, the billing platform, and the payer submission portal, replacing it demands organizational disruption that most buyers rationally avoid. And they expand within existing accounts. A company that starts with underwriting growing into claims, fraud detection, and pricing, each new module increasing the total switching cost of the platform.

> *The first vertical AI investment that really worked for me had nothing particularly impressive about the model underneath it. What it had was three years of domain-specific training data from real deployment, a team from inside the industry they were serving, and integrations so deep that the largest customer told me they'd rewritten their entire workflow around the product. That's a moat. The model is just the engine.*

> ***The moat isn't the model. The moat is what you build around it while everyone else is staring at it.***

💬 **TRACE'S SASSY SIDEBAR #4**

If I had a dollar for every pitch deck that said "proprietary AI" and then described a fine-tuned GPT wrapper, I could fund my own foundation model. Your model isn't proprietary. Your three years of dental claims data is.

Chapter 8

The Vertical AI Moat Stack

Layer by layer: how defensibility compounds from domain expertise to switching cost

Defensibility in vertical AI isn't a single advantage. It's a stack of overlapping advantages that compound on each other, making the total switching cost far larger than any individual layer would suggest.

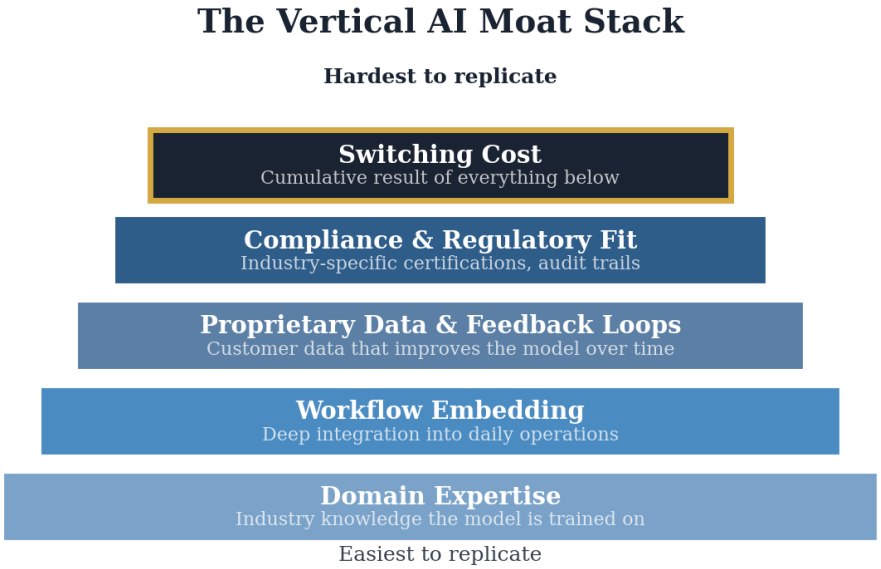

The Vertical AI Moat Stack. Each layer builds on the one below. Switching cost, at the top, is the cumulative result of everything beneath it.

Take a healthcare revenue cycle AI company. At the base: domain expertise—the team understands denial codes, payer behavior, and regulatory nuance a generalist AI would take years to learn. Layer two: workflow embedding, plugging directly into Epic or Cerner. Layer three: proprietary data—every claim processed generates labeled outcomes that improve the model. Layer four: HIPAA and SOC 2 compliance. Layer five: switching cost—moving to a competitor means retraining staff, recertifying systems, and losing years of model performance. Nobody switches. That's the moat.

When evaluating a vertical AI company, ask: how many layers are in place? One or two means vulnerable. All five means genuinely difficult to displace—and that difficulty is exactly what you want to own.

> ***In vertical AI, friction is the feature. High switching cost isn't a side effect; it's the goal.***

Chapter 9

Vertical AI Economics

Why 120% NRR changes everything, and how to read the numbers that matter

Strategy is only valuable if it shows up in the numbers. For vertical AI companies with embedded workflows and high switching cost, the key metric is net revenue retention: the percentage of revenue kept from an existing customer cohort after twelve months, including expansion and net of churn.

For a horizontal AI tool with real churn risk, NRR might be 85–95%. For a well-embedded vertical AI platform with expansion modules and genuine switching cost, NRR of 115–130% is achievable. That 30-point spread, compounded over five years, produces dramatically different businesses, not because of better salespeople, but because the product architecture creates natural expansion.

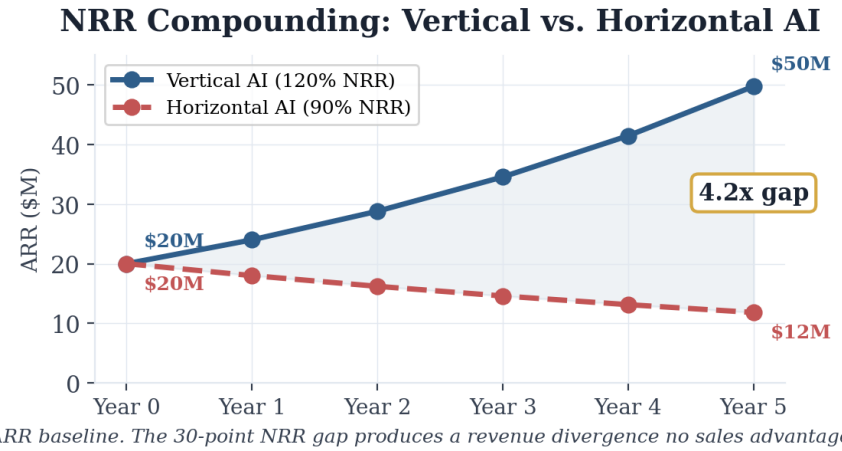

NRR compounding over five years from a $20M ARR baseline. The 30-point gap between horizontal and vertical AI produces a revenue divergence that no sales advantage can close.

Burn Multiple

Net cash burned ÷ net new ARR. Below 1x: elite. 1–2x: efficient. 2–3x: acceptable. Above 3x: concerning.

> ***When retention and switching cost are real, not just claimed in a deck, downside risk compresses. That matters more than people think when you're modeling exits under uncertainty.***

Chapter 10

Building for M&A vs. IPO

Most companies won't go public. Here's how to prepare for the exit that will actually happen.

Every founder I've met has, at some point, described their IPO. That ambition is real and it drives people to extraordinary things. It's also statistically uncommon. Planning capital strategy exclusively around an IPO is one of the most consistent ways founders and their investors set themselves up for friction.

Most venture-backed companies exit via acquisition. Strategic buyers are more active than public markets in almost every macro environment, and they evaluate acquisition targets on strategic logic that doesn't require the revenue scale thresholds public markets demand.

A company building for M&A focuses on integration compatibility: clean architecture, well-documented APIs, audited financials, and cultivated relationships with corp dev teams at likely acquirers. A company building for IPO needs $100M+ in ARR, governance maturity, multi-year revenue visibility, and SEC-ready compliance infrastructure. These requirements create real resource allocation tradeoffs. The founders best positioned are those who prepare for both and maintain optionality.

> **The Founder Who Prepared for Both**
>
> *A founder I know, I'll call him David, ran a vertical AI company in insurance with a clear IPO ambition. He also had a CFO who spent ten percent of her time maintaining relationships with corp dev teams at the five most logical strategic acquirers.*
>
> *In year four, one of those acquirers made an approach. The valuation was below his IPO target. But the integration was clean, diligence was fast because the financials were already audit-ready, and the transaction closed in four months.*
>
> *Optionality is the point. Prepare for the IPO. Maintain the relationships that make a great acquisition possible on your terms.*

> ***Let the market decide which path opens first. The founders who maintain optionality, and stay honest about which outcome is more likely, make the best decisions when the moment actually comes.***

❓ POP QUIZ — Chapter 10

1. Name three layers of the Vertical AI Moat Stack.
2. What NRR range distinguishes vertical AI from horizontal AI tools?
3. Why are foundation models often poor venture investments despite being technically impressive?

Answers (no peeking): 1. Domain expertise, workflow embedding, proprietary data, compliance infrastructure, switching cost. 2. Vertical AI: 115–130%. Horizontal AI: 85–95%. That 30-point gap compounds dramatically. 3. Training costs require sustained capital, competition is intense among well-funded players, and pricing pressure increases as compute costs fall.

💬 TRACE'S SASSY SIDEBAR #5

Here's a dirty secret: most VCs say they want IPOs. What they actually want is liquidity. If a $400M acquisition in year four returns 8x, that's a better story than a $2B IPO in year eleven where you've been diluted to 3%. Do the math before you do the dreaming.

📝 Notes & Scribbles

Think about an AI company you've seen, built, or invested in. How many layers of the Vertical AI Moat Stack does it actually have? Be honest. List them here.

P A R T III

The Capital Stack

Ownership, fund math, LP dynamics, and how to build a fund that actually works

Key Themes

The ownership geometry that makes small funds powerful

How LPs actually make allocation decisions

The six dimensions of institutional due diligence

Why raising too much capital destroys your edge

What You'll Learn

The ownership geometry that makes $75M funds mathematically superior to $750M funds for many strategies

How institutional LPs actually make allocation decisions—and why merit is only part of the equation

The six dimensions of LP due diligence that most first-time managers aren't prepared for

Why fund size drift is the silent killer of emerging manager franchises

How to model portfolio construction with precise numbers, not ranges

📝 Reader Challenge: If you're raising a fund, write down your target fund size. Now calculate: at your target ownership and the median exit distribution for your stage, how many companies need to exit above $500M for the math to work? If you're a founder, calculate what percentage of your company you'll own at a Series C.

Chapter 11

Emerging Managers: Ownership, Asymmetry, and Edge

The math is not just theoretical. According to Sapphire Partners' analysis of twenty years of fund data, only 50% of Fund I managers survive to raise a Fund II—an average of 3.4 years after first close. Just 17% reach Fund IV. Fewer than 5% ever reach Fund VIII, the threshold of a franchise firm. The median time from Fund I to franchise status is 17.7 years.

Yet the managers who survive—particularly those who raised during downturns—consistently outperform.

Cambridge Associates found that 40–70% of total venture gains over the past decade came from new and emerging managers. PitchBook's simulation of 2010–2019 vintages showed that a dollar allocated exclusively to emerging managers had a better probability of higher returns than one allocated to established funds, with the widest outperformance gap—a 4.5% median IRR advantage—coming from managers who raised during the post-GFC contraction.

The outperformance traces directly to the ownership geometry described above: smaller funds, higher ownership, better alignment with the actual exit distribution.

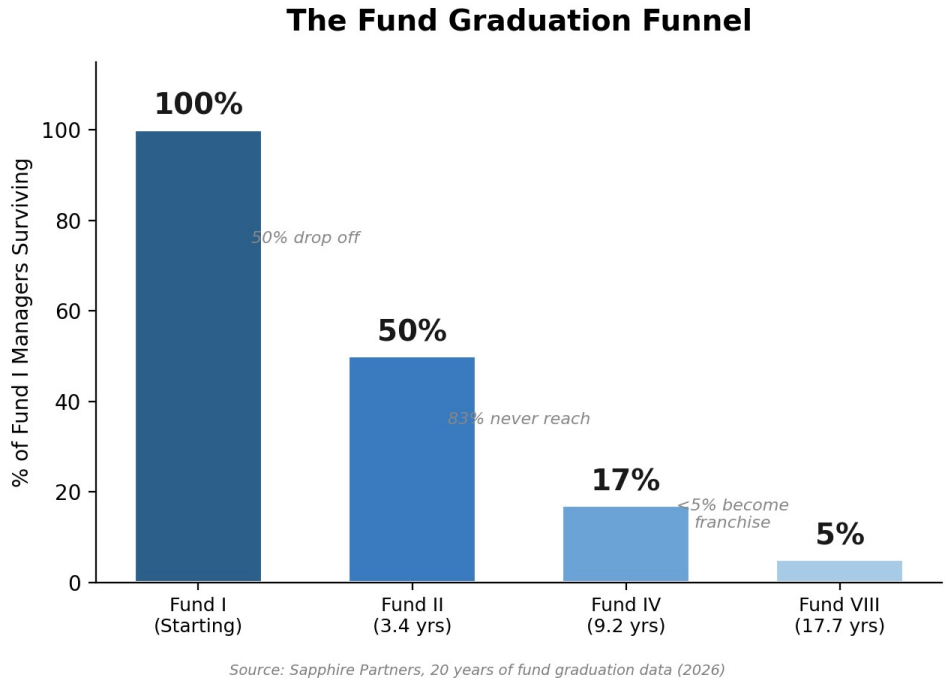

The fund graduation funnel. Half of all Fund I managers never raise again. Fewer than one in twenty become franchises. The timeline from launch to establishment is nearly a decade. Source: Sapphire Partners, 20 years of data.

The math that makes small funds powerful, and the mistakes that destroy that power

The venture industry has a conventional wisdom problem: it assumes bigger funds are better funds. More capital, more credibility, more support. More everything.

The math disagrees.

An emerging manager's true advantage isn't brand or hustle, even though those things matter. It's ownership geometry: the mathematical relationship between fund size, check size, ownership percentage, and the actual distribution of exits in the market.

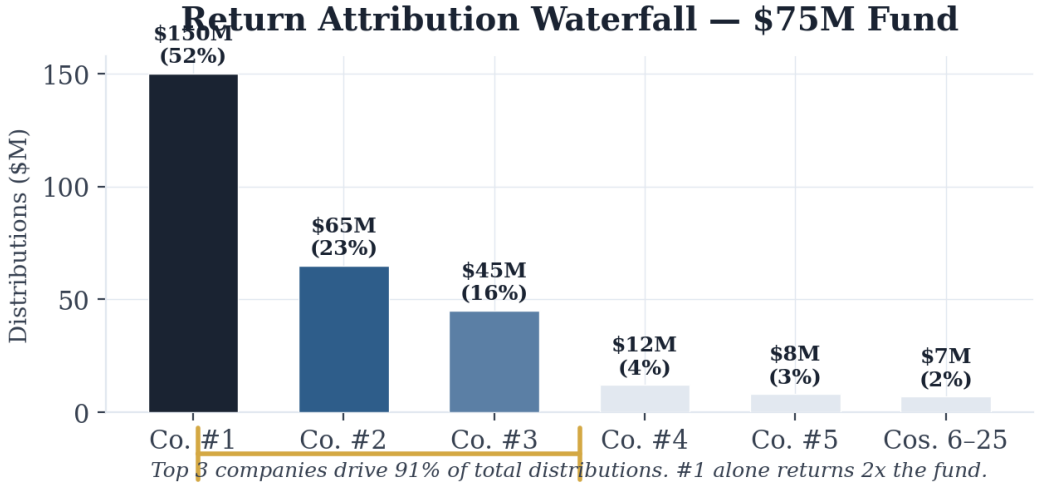

Return attribution waterfall for a $75M fund base case. The top three companies drive 78% of total distributions.

Here's the math made concrete. A $75M fund needs roughly $225M in distributions. With 10% ownership at exit, a $700M acquisition returns $70M. Two of those plus a handful of $200–500M exits and the fund returns 3x. Tight, but achievable—because the ownership math scales linearly when check sizes stay disciplined.

Now run it for a $750M fund. Same exits, same ownership. Those two companies generate $100M against a $2.25B required return—4.4% of the fund. You'd need twenty such outcomes to hit 3x. The math doesn't break from bad investing. It breaks from fund size.

> *I've watched managers build great first funds at $50–75M and then raise $200M for Fund II because LPs wanted to give them more. The strategy didn't change. The ownership didn't change. The exit distribution didn't change. But the required outcomes doubled. Some figured it out. Some are still explaining*

> *underperformance to their LPs.*

> ***Ownership is leverage. Leverage only works when you're selective about what you put it on.***

The emerging managers who succeed in the next decade will not be the ones who raised the most capital or built the biggest brand. They will be the ones who understood the math deeply enough to resist the pressure to drift—who kept their funds sized to the exit distribution, their ownership targets aligned with their return requirements, and their decision-making grounded in structural advantage rather than sentiment. The data now confirms what the math always implied: discipline compounds.

Chapter 12

How LPs Actually Think and How to Raise From Them

Power laws, pacing cycles, and the mechanics of a successful fund raise

Two things are simultaneously true about LP capital: it’s the lifeblood of every fund, and it behaves in ways that have almost nothing to do with the quality of any individual opportunity. Understanding both halves of that reality is what separates managers who raise efficiently from those who exhaust themselves in the process.

The Power Law Reality

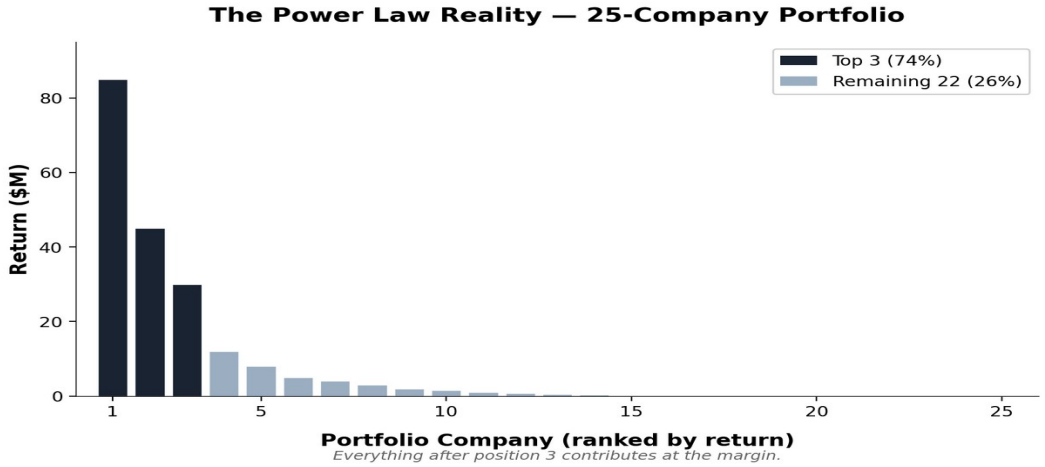

A simulated 25-company portfolio return distribution. The top three companies generate 74% of total fund returns. Everything else contributes at the margin.

A typical early-stage portfolio: 40–55% complete losses, 20–30% modest outcomes below 3x, 10–15% meaningful outcomes between 3–10x, and 2–5% extreme outliers above 20x. The fund’s performance depends almost entirely on that last category, sometimes 5% of the portfolio, sometimes one company.

This means reserves matter enormously. Without follow-on participation in your breakout companies, you dilute away your most valuable positions exactly when it costs the most. An investor who enters at seed with 12% and doesn't participate in Series A may end up with 4% by exit. One who exercises pro-rata stays near 9%. That three-point difference, across a $500M exit, is the difference between a $15M return and a $45M return.

The power law math is why fund construction matters so much. But construction is only relevant if you can raise the fund, and that requires understanding something most emerging managers underestimate: how LPs

actually make decisions, and why those decisions often have almost nothing to do with your portfolio companies.

How LPs Actually Decide

Institutional LPs allocate across asset classes with defined targets. Venture commitments exist within allocation frameworks that impose mechanical constraints on pacing regardless of how strong any individual opportunity looks.

Large platforms secure re-up commitments through relationship continuity. Emerging managers compete for genuinely discretionary commitments, the capital not already spoken for. In constrained markets, that pool shrinks. It isn't personal. It's portfolio mechanics.

The most effective emerging managers treat LP relationships as long-cycle investments, building trust and reporting cadence across years before any formal raise begins. They track LP pacing cycles the way traders track markets. They lead with intellectual rigor: a specific portfolio construction plan, ownership strategy expressed in precise numbers, exit modeling grounded in empirical distribution data rather than aspirational projections.

> ***Fundraising is not persuasion. It is alignment between your fund model and your allocator's constraints. The sooner you understand that, the more efficiently you'll raise.***

Chapter 13

LP Due Diligence: What Allocators Actually Evaluate

The six dimensions of LP diligence, and how to prepare for each one honestly

Most first-time fund managers prepare for LP meetings by polishing the investment thesis, rehearsing the team story, and updating portfolio slides. Those things matter. But sophisticated institutional LPs run a more systematic evaluation than most managers anticipate, and knowing the framework changes how you prepare.

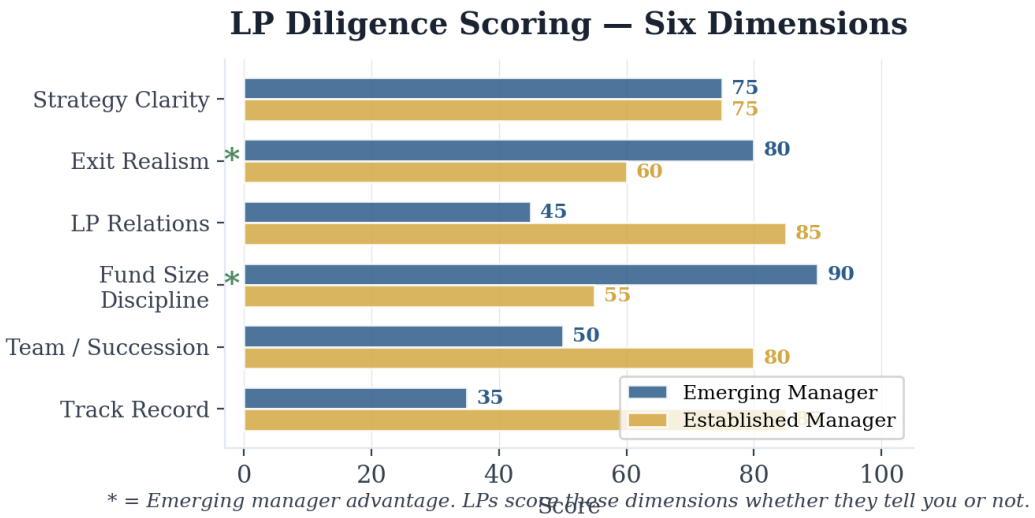

LP diligence scoring across six dimensions. Emerging managers have real advantages in fund size discipline and exit realism.

Sophisticated institutional LPs evaluate emerging managers across six distinct dimensions. Most first-time managers prepare for one or two. The ones who prepare for all six raise faster.

First, track record and attribution. LPs want to understand specific sourcing. Did you find this deal yourself, or did it come through your prior firm's platform? Is the edge portable to a new vehicle? Attribution matters as much as the return multiple.

Second, portfolio construction. Can you describe, in precise numbers, how your construction generates the return profile you're targeting? Initial check size, target ownership, reserve ratio, exit ownership assumption. Ranges signal that the plan isn't actually concrete.

Third, fund size discipline. Correctly sizing the fund to your deal access and exit distribution signals that you understand the math, not just the marketing. Overcapitalized relative to your strategy is a red flag, not a sign of confidence.

Fourth, team and succession. For solo GPs, proactively address key-person risk. How does the portfolio get managed if something happens to you? Have an answer before they ask.

Fifth, LP relations and reporting. References matter enormously. Managers who stayed transparent during difficult portfolio situations get materially better references than those who went dark when the news was bad. Proactively share your reporting templates.

Sixth, exit realism. Show that you've modeled the empirical exit distribution, not the aspirational one. Know the historical exit frequency in your specific stage and sector. Build your fund model around the median, not the tail.

> **The LP Meeting That Went Differently**
>
> *A manager I know was raising her first institutional fund with a strong track record but no realized exits. Her fund model showed three $1B+ exits from a 20-company portfolio.*
>
> *An endowment CIO stopped her mid-presentation and asked: 'What percentage of seed-stage healthcare software companies in the last decade achieved exits above $500M?' She didn't know.*
>
> *She went home, did the research (the answer was roughly 4%, or about one in twenty-five) and rebuilt her model around that number. More conservative. Also more credible.*
>
> *The same endowment committed to her fund six weeks later.*

> ***LPs don't fund optimism. They fund clarity. The clearer and more honest your model, the faster they move.***

Chapter 14

Fund Size Drift and Return Compression

The most common way emerging managers accidentally destroy their edge

Here is a pattern I've watched play out more times than I'd like. A manager closes a strong first fund. Early marks look good. LPs want to give them more capital. Fund II launches at two or three times the size of Fund I. The strategy doesn't change. The team doesn't change. The required outcomes double.

This is fund size drift, and it is the quiet killer of many emerging manager franchises.

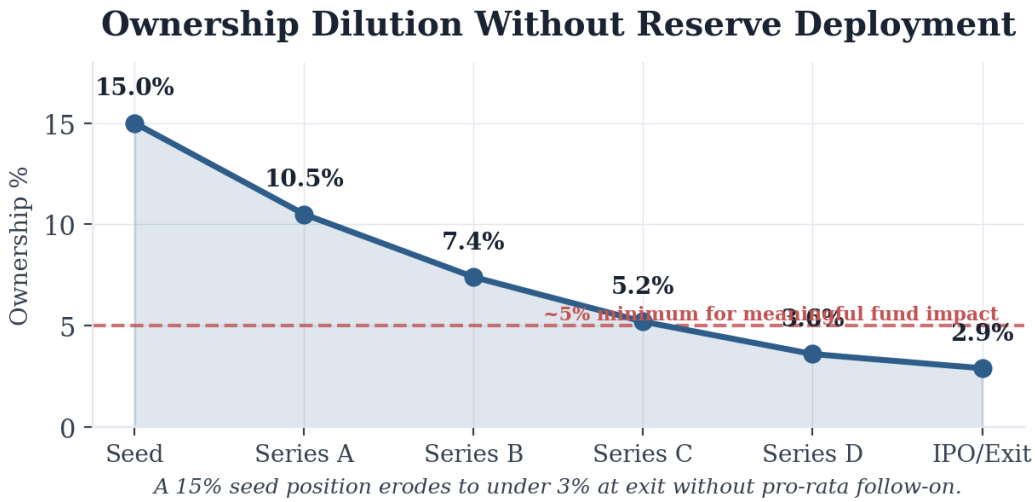

Ownership dilution across financing stages without active reserve deployment. It's easy to slip below the minimum ownership level required for any single exit to move the fund.

The adaptations that look rational in isolation are collectively destructive: writing larger checks (which requires more competitive rounds with compressed ownership), moving later stage (where early-stage advantages don't apply), and accepting lower ownership in hot deals. Each step erodes the ownership leverage that made the smaller fund work.

The right question before every fund raise is not 'how much will LPs give me?' It is 'what is the largest amount I can deploy while maintaining the ownership levels and deal stage my track record actually represents?'

> ***The optimal fund size is not the most you can raise. It's the most you can deploy at the ownership levels your strategy requires.***

❓ **POP QUIZ — Chapter 14**

1. What is “fund size drift” and why is it dangerous?
2. An LP targets 15% private allocation. Public markets drop 25%. What happens to their effective allocation?
3. What six dimensions do sophisticated LPs evaluate in diligence?

Answers (no peeking): 1. Raising Fund II at 2–3x Fund I without changing strategy. Required outcomes double but deal access doesn’t scale proportionally. 2. It mechanically rises to ~19%, pushing them over target and forcing them to pause new commitments. 3. Track record/attribution, portfolio construction, fund size discipline, team/succession, LP relations/reporting, exit realism.

📝 Notes & Scribbles

If you’re raising a fund: what is the maximum amount you can deploy at the ownership levels your strategy requires? If you’re a founder: what does your full dilution waterfall look like at exit? Sketch it here.

P A R T I V

The Operator’s Manual

Practical frameworks for founders and fund managers who are in the middle of it right now

Key Themes

How to treat fundraising as a pipeline, not a verdict

The Series A gap and how to prepare for it from day one

Burn discipline as a source of optionality

Capital stack sequencing and the compounding cost of early decisions

The alignment conversation nobody has before the term sheet

What You’ll Learn

Why treating fundraising as a probabilistic pipeline transforms your conversion rate and your mental health

The specific metrics Series A investors underwrite to—and how to start measuring them from day one

How to build a burn stress test that assumes everything takes twice as long as planned

Why every round you take creates obligations that compound through every subsequent round

The alignment conversation that prevents $25M from being left on the table

📝 Reader Challenge: Build your own fundraising stress test right now. Assume your next round takes 24 months instead of 12. Assume valuation multiples compress 40%. Assume your top customer delays renewal. Can you extend runway to 18 months without losing key people? Write down the honest answer.

Chapter 15

Fundraising Is a Funnel

Stop treating rejection as feedback. Start treating fundraising as pipeline.

Fundraising is deeply personal. When someone says no, it feels like a judgment. It often isn't.

The most useful reframe I've found, for founders raising rounds and managers raising funds alike, is to treat fundraising as a probabilistic pipeline, not a verdict process. Once you do that, early rejections stop feeling like judgments and start feeling like what they almost always are: stage-of-process mismatches, allocation timing issues, or portfolio construction constraints that have nothing to do with the quality of what you're building.

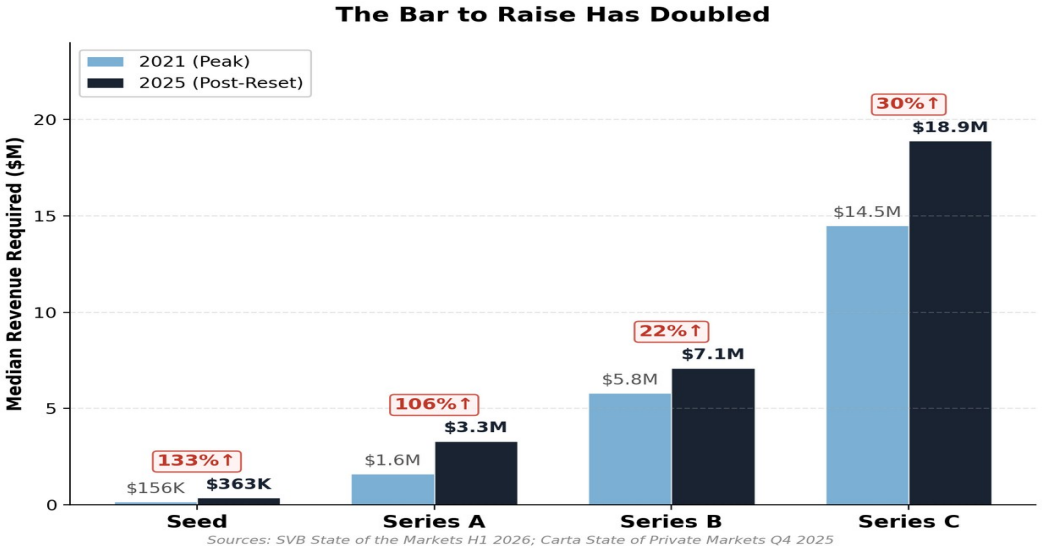

The fundraising funnel for a $75M emerging manager fund. The top of the funnel needs to be far wider than most managers expect.

The empirical funnel for an emerging manager raise: 150–200 LP conversations, 35–50 serious follow-ups, 15–20 in diligence, 8–12 commitments. A 2.4% close rate from initial conversation to commitment is typical for a first-time manager. That's the conversion rate.

> *When I was raising capital early in my career, I treated every rejection as a data point about my strategy. I'd revise the pitch, adjust the thesis, second-guess the focus. Eventually I realized the data point I was missing was simpler: most rejections had nothing to do with me. Once I understood the funnel, I stopped rewriting the deck after every no and started building the pipeline wider instead.*

> ***Start building the funnel earlier than feels necessary, wider than feels comfortable, and with more process than you think you need. The math only works if you give it enough surface area to operate on.***

> 💬 **TRACE'S SASSY SIDEBAR #6**
>
> *I keep a spreadsheet of every LP I've ever met. It has 847 rows. You know how many committed? Fourteen. That's a 1.7% conversion rate. And that's considered good. Welcome to fundraising.*

The bar has risen materially since 2021. According to SVB and Carta data, the median revenue required to raise a Series A more than doubled: from $1.6 million in 2021 to $3.3 million in 2025. At seed, the threshold jumped from $156,000 to $363,000. At Series B, from $5.8 million to $7.1 million. At Series C, from $14.5 million to $18.9 million. Down rounds have fallen below 14%—the lowest rate in three years—but only because the companies that couldn't meet these higher bars simply didn't raise at all. The reset didn't lower the bar. It raised the floor permanently.

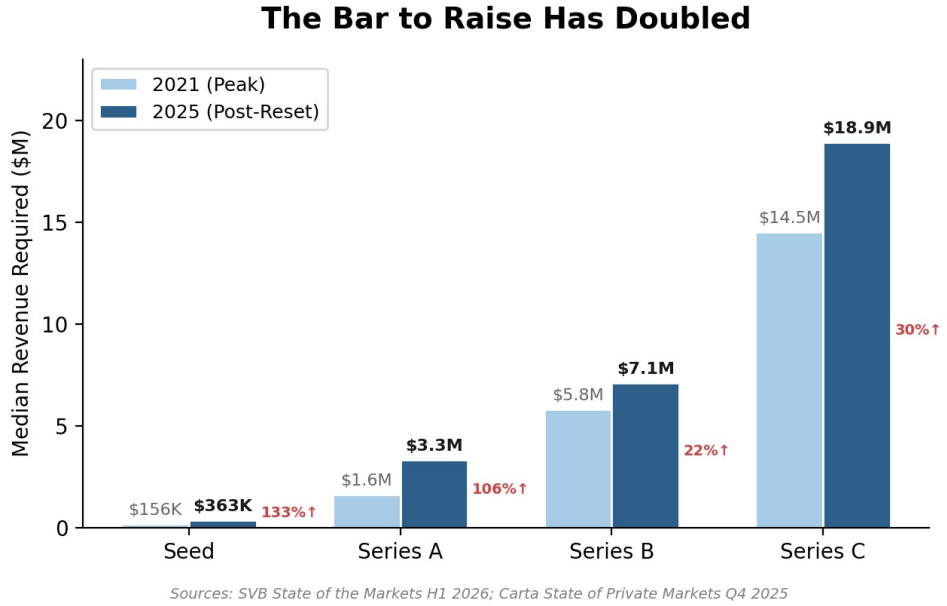

Median revenue required to raise, 2021 vs. 2025. The bar has roughly doubled at every stage. The companies that survived the reset earned their next round. Sources: SVB State of the Markets H1 2026; Carta State of Private Markets Q4 2025.

Chapter 16

The Series A Gap Is Real

Why seed capital doesn’t guarantee anything except eighteen months of runway

At some point between 2019 and 2022, the venture ecosystem collectively forgot that seed capital is not an outcome. It’s a starting line.

Seed rounds proliferated at extraordinary rates during the expansion. The bar lowered. Narrative mattered more, traction mattered less. Companies that would have struggled to raise $500K in 2018 closed $3M seed rounds in 2021. The capital flowed. Then the Series A market corrected.

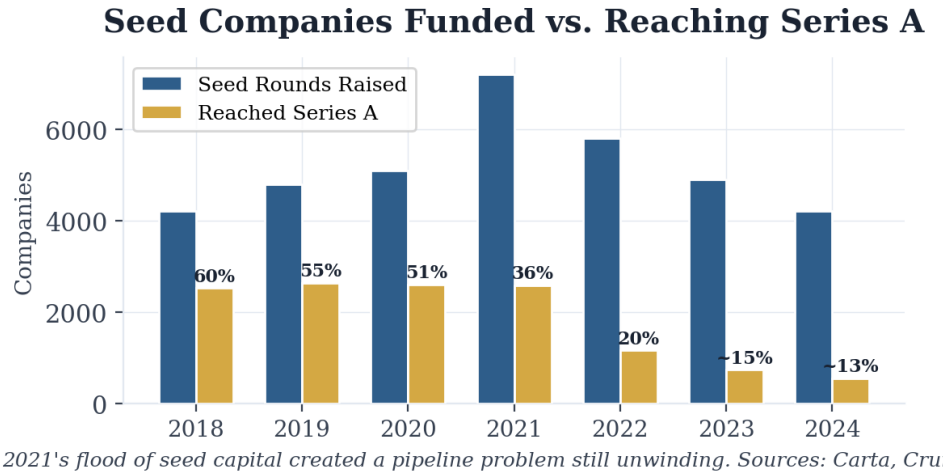

Seed companies funded vs. companies reaching Series A (2018–2024). The gap exists in every market environment.

Series A investors underwrite differently: $1–3M in ARR for SaaS, clear retention evidence, a validated acquisition channel, cohort data showing expansion. If the seed was funded on a story, the Series A is funded on evidence. Most stories don't become evidence in eighteen months.

I've had this conversation with founders shocked to be struggling at Series A. They'd raised pre-seed in weeks. Seed took a month. They assumed Series A would be a formality. It wasn't—because the evaluation framework changed completely. Narrative sufficiency no longer counts. Evidence sufficiency is what matters.

In one case, we rebuilt a founder's metrics dashboard over three months—tracking cohort retention, expansion revenue, and CAC with precision they'd never had. They raised their A six months later. The metrics hadn't changed. The ability to prove them had.

> ***Seed capital buys you eighteen months to prove revenue durability, retention quality, and unit economics. That's it. Spend it accordingly.***

💬 **TRACE'S SASSY SIDEBAR #7**

A founder once told me their Series A was "basically done" because three VCs said they were "super interested." I asked how many term sheets they had. Zero. Interested is not committed. Excited is not funded. Until the wire hits, you are still fundraising. Plan accordingly.

Chapter 17

Burn Discipline in a Repriced Market

Efficiency is not a constraint on ambition. It's what makes ambition sustainable.

Burn is not the enemy. Inefficient burn is the enemy.

Every venture-backed company burns cash by design. You're investing ahead of revenue to build something that can scale. The question is never whether you're burning. It's whether the burn is buying you something real: validated growth, durable retention, expanding accounts, defensible market position.

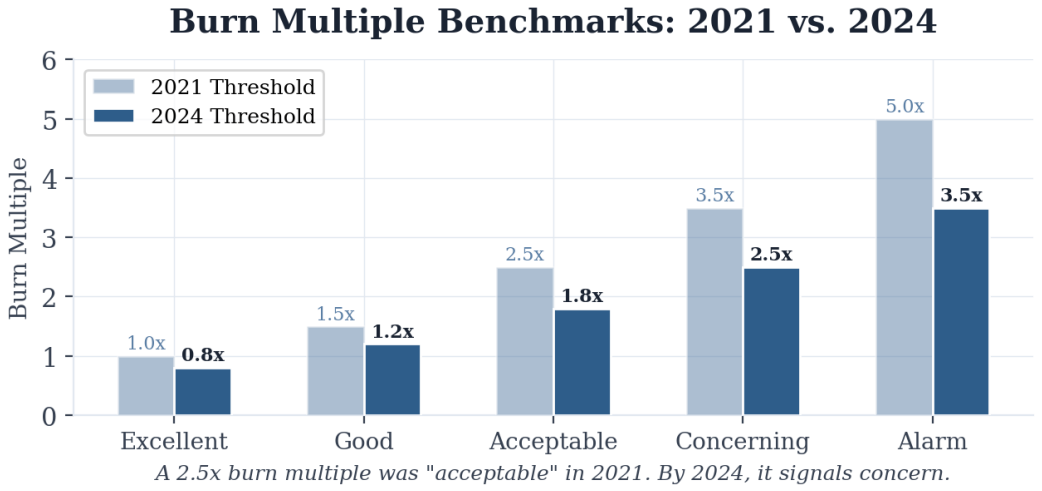

Burn multiple benchmarks and market median shifts from 2021 to 2024. The same burn multiple that was acceptable in an expansionary market signals concern today.

The most valuable exercise I recommend to every founder: build a stress test that assumes your next fundraise takes twenty-four months instead of twelve. Assume valuation multiples compress 40% from current marks. Assume one top account delays renewal. Can you extend runway to eighteen months without losing key people? If the answer is no, you don't have enough runway.

> ***Burn discipline creates optionality. Optionality prevents desperation. Desperation is the fastest path to a bad deal on someone else's terms.***

Chapter 18

Capital Stack Sequencing

Every round you take sets the terms for every round after it

Founders often think about financing as isolated events: the seed round is over, now we build, eventually we'll do the Series A. This framing is natural and almost entirely wrong.

Every round creates obligations, expectations, and governance that shape every subsequent round. The investors at seed will be at your board table for seven to ten years. The liquidation preferences you agree to will determine what gets paid in every exit scenario. The information rights you grant will determine who can call a board meeting at an inconvenient time.

The Capital Stack Across Financing Stages

Stage	Check Size	Valuation	Investors	Dilution
Pre-Seed	$250K - 1M$	$3-8M	Angels, Pre-Seed Funds	10-15%
Seed	$1-4M	$8-20M	Seed Funds, Micro VCs	15-20%
Series A	$5-15M	$20-60M	Early-Stage VCs	15-25%
Series B	$15-50M	$60-200M	Growth VCs	15-20%
Series C+	$50-200M	$200M-1B+	Late-Stage / Crossover	10-15%

Each stage carries different investor expectations, governance implications, and dilution patterns.

The capital stack across financing stages. Each stage carries different investor expectations, governance implications, and dilution patterns.

A founder entering seed at 80% drops to 55% post-seed, 38% after A, 26% after B, 18% at C. By exit, proceeds feel smaller than expected. Model the full waterfall before signing.

> *I've sat on both sides. As a founder I modeled the full dilution waterfall before signing. As an investor I reserved capital for pro-rata in my best companies. The ones where I exercised at Series A are where the math worked out.*

> ***Every round you take is a commitment, not just a transaction. Model the full stack before you sign the first term sheet.***

Chapter 19

Founder-Investor Alignment

The conversation no one wants to have before the term sheet, and everyone wishes they'd had

Of all the things that go wrong in venture-backed companies, misalignment between founders and investors is the most preventable and the most frequently ignored until it's too late.

Misalignment doesn't usually come from bad intentions. It comes from different mathematics. A seed investor who entered at a $5M valuation and owns 15% may achieve a strong fund return from a $150M acquisition. The founder, five years in and holding 40%, would receive $60M, a life-changing number. But if their investor is pushing for a $1B IPO exit, the conflict isn't about who's right. It's about who needs what from this outcome.

The Founder-Investor Alignment Matrix

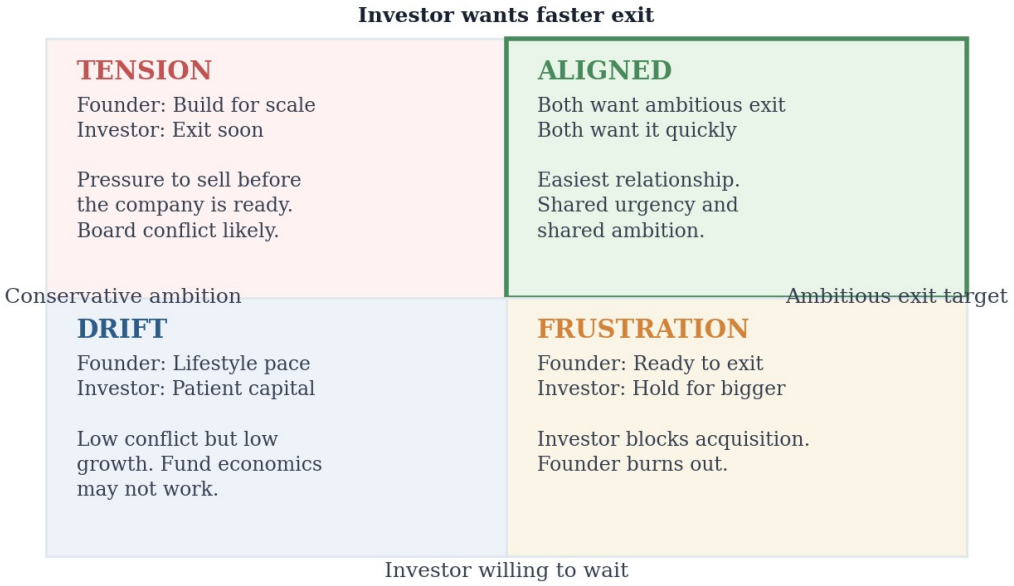

The misaligned quadrants are where good companies fall apart.

The founder-investor alignment matrix. The upper-right quadrant, both parties want an ambitious exit quickly, is where the relationship is easiest. The misaligned quadrants are where good companies fall apart.

The governance reality: liquidation preferences, protective provisions, information rights—all negotiated at signing—play out differently depending on alignment. A well-aligned investor with strong governance rights is an asset. A misaligned one with the same rights is a structural risk.

The Acquisition Offer

A founder I know, I'll call him Michael, received an $80M acquisition offer in year four. His seed investors, who'd entered at a $6M valuation, would see a 10x return. Michael, holding 42%, would receive $33.6M.

Michael wanted to take the offer. His Series A investor, who'd entered at a $40M valuation, needed a $200M+ exit to generate a meaningful return. They pushed hard against the deal.

The conflict wasn't about the company's prospects. It was about different entry prices creating different required outcomes. Neither party had ever discussed this explicitly before the offer arrived.

The deal fell apart. Two years later, in a tighter market, Michael sold for $55M. His Series A investor still didn't make meaningful money. Michael left $25M on the table from the earlier offer.

Have the conversation before the term sheet.

Here's something most VCs won't admit publicly but every founder deserves to know: we aren't value-add all the time. Our incentives shift with fund size, timelines shift with vintage, and board behavior shifts with portfolio context. Don't just diligence our track record. Diligence our alignment.

> ***Ask your investors what they need from this outcome before you take their money. The answer will tell you everything.***

❓ POP QUIZ — Chapter 19

1. A seed investor enters at $5M. A founder holds 40%. A $150M acquisition offer arrives. Why might the Series A investor push against the deal?
2. What is the single most consequential lever an emerging manager controls?
3. Why should founders model the full dilution waterfall before signing any term sheet?

Answers (no peeking): 1. The Series A investor entered at a higher valuation and needs a larger exit to generate meaningful returns. 2. Exit ownership percentage. A 2-point reduction (9% to 7%) reduces gross portfolio value by 22%. 3. Each round creates obligations that shape every subsequent round. Ownership at exit may be far lower than founders expect.

💬 TRACE'S SASSY SIDEBAR #7

The most awkward conversation in venture capital is also the most important one: "What number makes this work for you?" I've seen $25M left on the table because nobody asked that question before the term sheet. Ask it. Ask it at dinner. Ask it in the first meeting. Ask it before you like each other too much to be honest.

📝 Notes & Scribbles

What is your unfair advantage? If you can’t fill this space, that’s your answer. Be specific. “Hard work” isn’t an unfair advantage. What can you do that no one else can?

P A R T V

Structural Advantage

What endures across cycles, and how to position yourself to be there when it matters

Key Themes

Why AI makes experimentation cheap but differentiation expensive

How to read the market cycle and act accordingly

The discipline that compounds longer than sentiment

What You'll Learn

Why AI makes experimentation cheap but differentiation expensive—and where returns live in that gap

How to identify which phase of the market cycle you're in and which decisions belong to each

The chain of causation from fund size to ownership to exit distribution to returns

Why discipline endures longer than exuberance and structure compounds longer than sentiment

📝 Reader Challenge: Think about the last major investment decision you made—as a founder, investor, or employee. Was it driven by the current market sentiment or by structural analysis? Be honest. Write down what you'd do differently with the frameworks from this book.

Chapter 20

Capital Efficiency in the AI Era

Why AI makes experimentation cheap and differentiation expensive

AI has lowered the marginal cost of building software. A function that previously required ten engineers for six months can now be prototyped by two engineers in three weeks. This compression is real and has real implications for startup formation and competitive dynamics.

What it has not changed is the cost of building a durable company. Enterprise sales cycles haven't compressed. Regulatory compliance hasn't been automated away. The human trust required to become embedded in critical workflows still takes years to earn. The difference between a prototype and a production system that a Fortune 500 will bet its operations on is still enormous.

The opportunity in the AI era is not that building is cheap. It's that the gap between 'cheap to experiment' and 'expensive to differentiate' has widened, and that gap is where venture returns live. As more companies can build more quickly, competitive density in any application category increases rapidly. The advantage that early AI-native startups derived from technical capability compresses as those capabilities become accessible to everyone.

For a visual breakdown of how the J-curve plays out across a full fund lifecycle, including how DPI lags IRR in early years and why AI-era fund timelines are not materially shorter, see Appendix B.

> ***The AI era didn't change the fundamentals of venture. It amplified them. Durable companies are more durable. Fragile companies are more fragile. Know which one you're building or backing.***

Chapter 21

Reading the Market Cycle

How to know which phase you're in, and which decisions belong to each

Most venture practitioners spend their careers in one or two full market cycles. The compressed nature of fund lifecycles means a manager's track record may reflect primarily one macro environment rather than full-cycle performance. Understanding cycle dynamics, where you are, how long each phase lasts, what decisions each phase calls for, is one of the most valuable and least-taught skills in the industry.

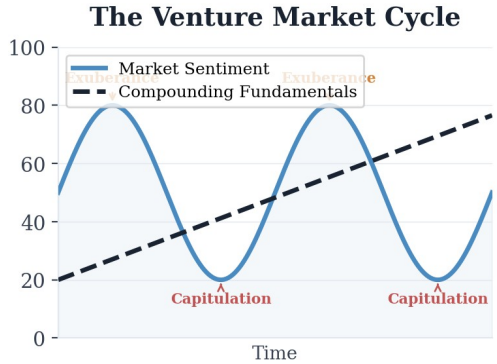

Sentiment oscillates; fundamentals compound. Ownership, exit distribution, and LP allocation operate through every phase.

The venture market cycle. Sentiment oscillates; fundamentals compound. The mechanics of ownership, exit distribution, and LP allocation operate through every phase.

Expansion: What It Hides

In expansion phases, nearly every manager looks good. Marks rise, capital is abundant, and the feedback loop between investment activity and valuation appreciation obscures a lot. The danger isn't that expansions last too long; it's that they make discipline seem optional. The managers who emerge from expansions intact are those who held ownership discipline when competitive pressure pushed them to accept lower stakes, and who reserved capital for breakout signals rather than averaging down.

Contraction: What It Reveals

Contractions are clarifying. They reveal which companies have genuine revenue durability and which were growing on loose underwriting and abundant capital. They reveal which fund models were built for the real exit distribution and which relied on optimistic assumptions that evaporate when the IPO window closes. For investors with dry powder, contractions represent some of the best vintage-year opportunities: entry valuations reflect constrained sentiment, competition decreases, and companies that survive

with quality intact are often dramatically better positioned than their marks suggest.

The current contraction is now visible in the data. According to Sapphire Partners, 2025 marked the first year since the dot-com collapse that the venture industry experienced a net decrease in active firms. More than 10,000 venture firms exist on paper in the U.S., per PitchBook, but fewer than half are actively investing. The rest are zombies—carrying unrealized marks, collecting management fees on legacy funds, but functionally absent from the market. New fund formation has fallen below the replacement rate. This is what a contraction looks like from inside: not a dramatic collapse, but a quiet thinning of the field that rewards the disciplined and exposes the over-extended. PitchBook's 2026 outlook called the difficulty for emerging managers to raise new funds the 'biggest risk in early-stage investing.' They're right—and it's also the biggest opportunity for the managers who do.

> **2021 and Its Lessons**
>
> *In 2021, global venture funding reached $621 billion. Seed valuations hit multiples previously reserved for growth-stage companies. Some managers who had built disciplined processes abandoned them under competitive pressure, accepting lower ownership, writing bigger checks into later stage, raising larger funds than their strategy warranted.*
>
> *By 2023, many of those companies faced down rounds. Many of those funds were explaining underperformance to LPs.*
>
> *The managers who navigated the period well had done two things consistently: held ownership discipline even when it meant losing deals, and built reserve positions in portfolio companies showing genuine revenue durability.*

The 2021 vintage included many excellent companies. It also included capital misallocated under the assumption that the cycle would continue indefinitely. It didn't. It never does.

> ***Expand your pipeline in contractions. Tighten your ownership discipline in expansions. The cycle tempts you to do the opposite of both.***

Chapter 22

The Discipline of Structure

What endures, and why it's the only thing worth building toward

Venture capital oscillates between exuberance and skepticism in ways that feel significant in the moment and predictable in retrospect. Structure operates beneath both.

Fund size determines required outcomes regardless of market sentiment. Ownership determines leverage regardless of prevailing multiple assumptions. Exit distributions constrain what's possible regardless of the narrative in any given year. These mechanics don't pause during optimism and they don't relax during pessimism. They compound continuously.

The repricing that began in 2022 didn't end venture capital. It clarified it. What survived the clarification, what emerged with performance intact, were funds with disciplined ownership, founders with durable retention, businesses with genuine switching cost, and investors who had modeled the real exit distribution rather than the aspirational one.

3x–5x
Net return target for top-quartile early-stage funds

8–12 yrs
Typical fund lifecycle from first investment to final distribution

The chain that determines venture outcomes is simple: asymmetric returns require concentrated ownership. Concentrated ownership requires disciplined check sizing. Disciplined check sizing requires fund size aligned with accessible exit distribution. Exit distribution alignment requires honest underwriting rather than aspirational modeling. That chain isn't exciting. It doesn't generate the headlines that venture mythology celebrates. But it is the foundation on which durable performance is built.

> ***Discipline endures longer than exuberance. Structure compounds longer than sentiment. The window that matters is the one you prepared for, not the one you hoped for.***

The opportunity in this era, the AI era, the repriced era, whatever name we eventually give it, is not to wait for expansion to return. It is to operate with clarity in whatever environment exists.

Because constraint is not an obstacle. It is a filter.

And what survives filters tends to endure.

❓ **POP QUIZ — Chapter 22**

1. What is the “narrative gap” in venture?
2. Expand your pipeline in ___. Tighten your ownership discipline in ___.
3. In one sentence: what endures longer than exuberance?

Answers (no peeking): 1. The distance between LP meeting projections and what the distribution history actually delivers. 2. Contractions; Expansions. The cycle tempts you to do the opposite. 3. Discipline. Structure compounds longer than sentiment.

📝 Notes & Scribbles

The window that matters is the one you prepared for. What are you preparing for right now? What decision are you avoiding? Write both down. The second one is usually more important.

Still In It

You now understand more about how venture capital works than most people who have been operating inside it for years.

That doesn't make it easier. The rules will still change. The goalposts will still move. Most startups won't make it. Most funds won't return what they promised. Most people who try, as founders, as investors, as operators, will face more rejection and uncertainty than they expected when they started.

That's not a reason not to try.

The founders who succeed aren't the ones who avoided the hard parts. They're the ones who understood what they were walking into. Who had a clear picture of the odds and showed up anyway. Who knew the math and still believed in the mission. Who understood that once you raise, you're always raising, and decided that was a price worth paying for the chance to build something real.

Here's what I've learned after fifteen years on both sides of the table: the bet we make at the earliest stage, before product, before revenue, before any of the metrics in this book exist, is always a bet on a person. Do you have the unfair advantage? The skills nobody else in this space has? The network that opens the right doors? The specific obsession that makes you the one who figures this out?

That's what this whole industry is, underneath all the fund mechanics and ownership geometry and exit distribution curves. It's a bet on human beings who see something others don't and refuse to stop building until the world catches up.

> ***It's supposed to be this hard. Only the ones who understand that, and show up anyway, were ever going to make it.***

If you read all of this and you're still in it, still building, still investing, still trying to figure out where you fit, then you were always the bet worth making.

Keep building.

Trace Cohen
New York, 2026

🔧 Interactive Tools

The frameworks in this book are meant to be used, not just read. Visit valueaddvc.com to access the interactive versions of the tools discussed in these chapters.

TOOL	WHAT IT DOES	CHAPTERS
Fund Economics Calculator	Model net returns across fund sizes, fee structures, and carry scenarios. Plug in your own numbers.	Ch. 3, 11, 14
Deal Memo Generator	AI-assisted deal memo creation with structured diligence templates and scoring frameworks.	Ch. 6, 7, 8, 13
Dilution Waterfall Simulator	Visualize founder and investor ownership from seed through Series C and exit scenarios.	Ch. 14, 18, 19
Burn Multiple Tracker	Track burn efficiency over time with benchmarks for your stage and sector.	Ch. 9, 17
LP Fundraising Pipeline	CRM-style tracking for LP conversations, follow-ups, and commitment pacing.	Ch. 12, 15
Exit Distribution Analyzer	Model portfolio returns under bear, base, and bull scenarios with real exit data.	Ch. 4, 11, App. B
Vertical AI Moat Scorecard	Score any vertical AI company across the five moat layers. Share with your team.	Ch. 7, 8
NRR Compounding Model	See how the 30-point NRR gap between vertical and horizontal compounds over 5 years.	Ch. 9

➡ valueaddvc.com

All tools are free for readers. Built with the same AI stack discussed in Chapter 2. If something breaks, that's also part of the lesson.

Afterword: The Data Has Arrived

When I started writing this book, I believed the structural arguments would prove themselves over time. Fund size determines probability. Ownership determines leverage. Exit distributions constrain what's possible regardless of sentiment. Cycles reward discipline and punish drift. These weren't predictions. They were mechanics.

By the time this book reached your hands, the data arrived faster than I expected.

In January 2026, Sapphire Partners published an analysis titled 'I See Dead VCs.' Their finding: 2025 was the first year since the dot-com collapse that the U.S. venture industry experienced a net decrease in active firms. More than 10,000 firms exist on paper. Fewer than half are actively investing. Their twenty-year dataset on fund graduation rates revealed that only 50% of Fund I managers survive to Fund II, just 17% reach Fund IV, and fewer than 5% ever reach Fund VIII. The median path from launch to franchise takes 17.7 years. In that same period, first-time funds raised a combined $6.6 billion in 2025. Andreessen Horowitz alone raised $15 billion in a single close.

Silicon Valley Bank's State of the Markets report showed $340 billion invested in U.S. venture-backed companies in 2025—the second-strongest year on record. But 47% of that capital flowed through deals of $500 million or more. Twenty-four companies received billion-dollar venture deals in a single year. Five AI companies captured $192 billion in cumulative funding, outvaluing every dot-com era IPO combined. The question SVB raised directly: 'Is a $1 billion venture capital deal really venture?' The bifurcation described in Chapter 3 is no longer theoretical. It is empirical.

Carta's State of Private Markets data confirmed that the revenue bar to raise has roughly doubled at every stage since 2021. Series A now requires a median of $3.3 million in revenue, up from $1.6 million. Down rounds fell below 14%—not because the market softened, but because the companies that couldn't clear the higher bar simply stopped raising. AI startups command a 38% valuation premium at Series A and 193% at Series E and beyond. The market isn't broken. It's bifurcated.

Cambridge Associates found that 40–70% of total venture gains over the past decade came from new and emerging managers. PitchBook's simulation showed that emerging managers outperformed established funds across 2010–2019 vintages, with the widest gap—a 4.5% median IRR advantage—appearing among managers who raised during the post-GFC freeze. Greenspring Associates documented a +0.87% net IRR advantage for emerging managers across 180 partnerships. The emerging manager advantage described in Chapter 11 is not a theory. It is a documented pattern with decades of data behind it.

Every structural argument in this book—fund size discipline, ownership geometry, the power law, the cyclical advantage of constraint—is now being validated by the most comprehensive LP and market data the industry has ever published. The mechanics didn't change. The data just caught up.

If you're reading this during the contraction, you're reading it at exactly the right time. The managers who build with discipline during this phase—who size their funds correctly, maintain ownership, and operate with clarity rather than sentiment—will define the next decade of venture returns. The data says so. The math always did.

Sources and Further Reading

The data cited in this afterword and throughout the updated chapters draws from the following reports, each of which is recommended reading for anyone serious about understanding the current venture landscape:

Sapphire Partners — "I See Dead VCs" (January 2026). Twenty years of fund graduation data, GP lifecycle analysis, and the first documented net decrease in active U.S. venture firms since 2001. sapphireventures.com

Silicon Valley Bank — State of the Markets H1 2026. Comprehensive analysis of deal size concentration, mega-deal trends, and the $192B AI capital stack. svb.com/trends

Carta — State of Private Markets Q4 2025. Revenue benchmarks by stage, valuation premiums for AI, and down-round frequency data. carta.com/blog

Cambridge Associates — Emerging manager performance analysis. Documents that 40–70% of total venture gains over the past decade came from new and emerging managers. cambridgeassociates.com

PitchBook / NVCA — Venture Monitor and Emerging Manager reports. Emerging manager simulation, first-time fund data ($6.6B in 2025), and the "biggest risk in early-stage investing" outlook. pitchbook.com

Greenspring Associates / StepStone — "The Case for Emerging Managers." Documented +0.87% net IRR advantage for emerging managers across 180 partnerships. stepstonegroup.com

Appendix A

Venture Capital Glossary and Metrics Reference

A plain-English reference for every term and metric used in this book. Benchmark ranges should be adapted based on stage, sector, and market conditions.

Fund-Level Metrics

DPI	Cash returned to LPs ÷ invested capital. The only metric that counts cash in the bank. Top-quartile funds: 2.0x+ DPI at close.
TVPI	DPI plus remaining portfolio value ÷ invested capital. Useful mid-fund but subject to mark optimism. Top-quartile 5-year: 1.8x+.
Net IRR	Time-weighted return accounting for timing of cash flows. Top-quartile early stage: 25%+. Early returns worth more than equivalent late returns.
Management Fee	Typically 2%/year during investment period, stepping to 1.5%/year during harvest. On a $75M fund: ~$13–15M total over ten years.
Carried Interest	20% of profits above an 8% IRR hurdle. The primary mechanism aligning GP and LP incentives.

Company-Level Metrics

NRR	Revenue kept from existing customers after 12 months including expansion. Above 100% = growth without new customers. Elite vertical AI: 115–130%.
Burn Multiple	Net cash burned ÷ net new ARR. Below 1x: elite. 1–2x: efficient. 2–3x: acceptable. Above 3x: concerning.
ARR	Annualized value of recurring contracts. Series A benchmark: $1–3M ARR typical for SaaS.
Gross Margin	Revenue minus COGS ÷ revenue. Vertical AI SaaS target: 65–80%. Below 50% suggests infrastructure or services cost issues.
Rule of 40	Revenue growth rate + EBITDA margin. Sum above 40% considered healthy at growth stage and above.

Appendix A (continued)

Venture Glossary

Additional terms used throughout this book, in plain English.

Anchor Investor	The first LP to commit to a fund raise, typically at 10–25% of the target, lending credibility to attract the rest.
Anti-Dilution	A provision protecting investors from the full impact of a down round. Weighted-average is standard and founder-friendly; full-ratchet is punitive and rare.
Bridge Round	Short-term financing, often convertible, used to extend runway between priced rounds.
Cap Table	The full breakdown of who owns what: founders, employees (via options), and all investors.
Carry / Carried Interest	The fund manager's share of profits, typically 20% above the hurdle rate.
Convertible Note	Debt that converts to equity at a future priced round, typically at a discount or with a valuation cap.
D/D (Due Diligence)	The research process investors use before making an investment. Financial, legal, technical, and reference-based.
Denominator Effect	When public market values decline, private allocations rise mechanically as a percentage of total portfolio, forcing LP rebalancing.
Down Round	A financing round at a valuation below the prior round.
DPI	Distributed to Paid-In. Cash returned to LPs as a multiple of invested capital.
Drag-Along Rights	Majority shareholders' right to compel minority shareholders to approve a sale.
Dry Powder	Capital committed but not yet deployed.
IRR	Internal Rate of Return. Time-weighted return accounting for timing of all cash flows.
J-Curve	The pattern of negative early returns (from fees and write-offs) followed by positive returns as exits materialize.
Lead Investor	The investor who sets the terms, takes the largest check, and typically joins the board.
Liquidation Preference	The right to receive a specified return before common shareholders in an exit. 1x non-participating is standard.

MOIC	Multiple on Invested Capital. How many times you got back what you put in.
Pro-Rata Rights	The right to invest in future rounds proportionally, to maintain ownership percentage.
SAFE	Simple Agreement for Future Equity. Converts to equity at a future priced round, without debt structure.
TVPI	Total Value to Paid-In. DPI plus current estimated value of unrealized investments.
Unicorn	A private venture-backed company valued at $1 billion or more.
Vintage Year	The year a fund makes its first investment, used for peer comparison.
Waterfall	The order of distributions in an exit: creditors first, then preferred per liquidation preferences, then common.

😏 The Honest Glossary

What people in venture actually mean when they say these things. Not for the easily offended. (If you are easily offended, you may be in the wrong industry.)

“Value add”	*I’ll introduce you to three people, two of whom won’t respond. The third will take a meeting and then ghost you. You’re welcome.*
“Strategic partnership”	*We’re dating but not exclusive. If a better deal comes along, we’ll call it “portfolio optimization.”*
“We’re very excited about your company”	*We are going to pass. The enthusiasm is inversely correlated with commitment.*
“Let’s keep in touch”	*I am never going to contact you again unless you make the front page of TechCrunch.*
“We’re thesis-driven”	*We invest in whatever our best-performing peer invested in last quarter, but with a narrative.*
“Proprietary deal flow”	*Our associate found them on LinkedIn. Like everyone else.*
“We have strong conviction”	*We followed the lead investor and then added ourselves to the press release.*
“The market will tell us”	*We don’t have an opinion and we’re not going to form one until someone braver goes first.*
“We’re founder-friendly”	*Until the first down round, at which point we will become “governance-focused.”*
“We bring operational expertise”	*One of our partners once worked at a startup for eleven months in 2014.*
“Emerging manager premium”	*You’re too small for us but we don’t want to say that, so here’s a fancy word for it.*

“We need to see more traction”	*Come back when someone else has already validated you so we don’t have to take actual risk.*
“We love the team”	*We don’t understand the product.*
“Interesting space”	*I have never heard of this before and I am Googling it while you talk.*
“We’re looking at the space”	*We made one investment that accidentally touched this sector and now we claim expertise.*
“We want to be helpful, not just a check”	*We are definitely just a check. But a very enthusiastic one.*
“Our LP base is very sophisticated”	*Our largest LP is my uncle.*
“We have a data-driven approach”	*We made a spreadsheet once.*

If you recognized yourself in more than three of these, congratulations: you are in venture capital. If you recognized yourself in all of them, you might be a managing partner.

Appendix B

Scenario Models: Bear, Base, and Bull Case

Three scenarios for a $75M emerging manager fund. Illustrative, based on realistic historical data, designed to show how mechanics interact.

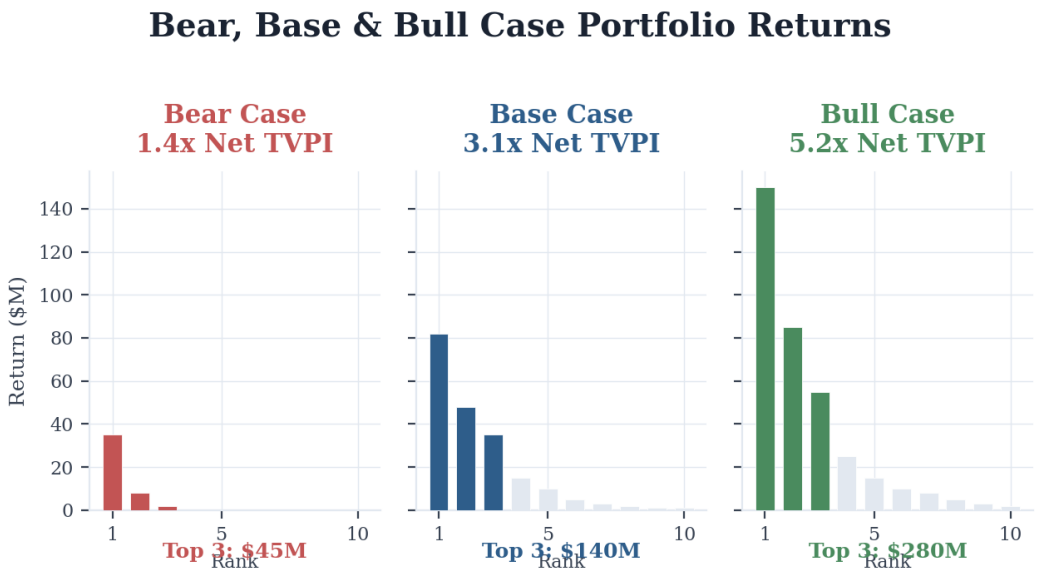

Bear, Base, and Bull case portfolio return distributions. The entire difference between 1.4x and 5.2x net returns is driven by whether one or two breakout companies materialize in the top three positions.

Fund assumptions (all scenarios): $75M fund. Management fee 2%/year investment period, 1.5%/year harvest. Total fees ~$13.5M. Carried interest 20% above 8% IRR hurdle. 22 portfolio companies. Average initial check $2.5M. Average initial ownership 14%. Average exit ownership (post-dilution) 9%.

Bear Case: 1.4x Net MOIC

No breakout exits. Top outcome: $180M acquisition at 9% ($16M). Three exits at $60–120M contributing $18M. Remainder: $12M. Gross portfolio: $46M. Net to LPs: ~$105M, or 1.4x.

Base Case: 2.8x Net MOIC

One company exits at $600M at 9% ($54M). Second: $350M at 9% ($31M). Third: $200M at 10% ($20M). Four additional contribute $26M. Write-offs:

45% by number. Net to LPs: ~$210M, or 2.8x. Achievable with disciplined construction and no exceptional luck required.

Bull Case: 5.2x Net MOIC

One company: $1.5B IPO at 8% ($120M). Second: $750M at 9% ($67M). Third: $400M at 10% ($40M). Six additional contribute $65M. Net to LPs: ~$390M, or 5.2x.

Key sensitivity: A 2-percentage-point reduction in average exit ownership (from 9% to 7%) reduces gross portfolio value by 22% across all scenarios. This single variable, managed through initial ownership discipline and pro-rata exercise, is the most consequential lever available to an emerging manager.

Appendix C

Suggested Reading

The books that shaped the thinking in these pages.

On Venture Capital

Venture Deals by Brad Feld & Jason Mendelson. *Read it before you sign anything.*

The Power Law by Sebastian Mallaby. *The best history of venture capital written.*

Secrets of Sand Hill Road by Scott Kupor. *How venture funds actually operate.*

Zero to One by Peter Thiel. *The canonical case for building monopolies over competing in existing markets.*

On Strategy and Competitive Dynamics

Competitive Strategy by Michael Porter. *Still the best starting point for understanding industry structure.*

The Innovator's Dilemma by Clayton Christensen. *How incumbents fail to respond to disruption.*

7 Powers by Hamilton Helmer. *The modern framework for competitive moats.*

Good Strategy / Bad Strategy by Richard Rumelt. *Clarifies what strategy actually is vs. what it's confused with.*

On Capital Markets and Risk

Capital Returns by Edward Chancellor (ed.). *The best framework for understanding how capital allocation affects returns over time.*

The Intelligent Investor by Benjamin Graham. *Margin of safety translates directly to private market discipline.*

Against the Gods by Peter Bernstein. *Essential grounding for probabilistic thinking.*

On AI and Technology

The Coming Wave by Mustafa Suleyman. *A grounded examination of where AI development is heading.*

AI Superpowers by Kai-Fu Lee. *Useful framework for thinking about AI commercialization across markets.*

Index

Key concepts indexed by chapter. Full page numbers available in the digital edition.

A

AI stack (layers) Ch. 6 | AI, vertical vs. horizontal Ch. 6, 7, 8, 9 | ARR Ch. 9, 16, App. A | Anchor investor Ch. 12, App. B | Anti-dilution Ch. 18, App. B

B

Bifurcation (fund sizes) Ch. 3 | Board composition Ch. 19 | Burn multiple Ch. 9, 17, App. A | Building with AI Ch. 2

C

Capital cost and rates Ch. 1, Intro | Capital stack sequencing Ch. 18 | Carry App. A, B | Cap table App. B | Case studies Ch. 4, 13, 14, 19, 21

D

D/D (Due Diligence) Ch. 13, App. B | Data moats Ch. 7, 8 | Denominator effect Ch. 1, 5 | Dilution Ch. 11, 14, 18 | DPI Ch. 5, 12, App. A | Down round Ch. 1, App. B

E–F

Emerging managers Ch. 3, 11 | Exit distribution Ch. 4, 11, App. C | Exit planning (M&A vs. IPO) Ch. 10 | Foundation models Ch. 6 | Founder-investor alignment Ch. 19 | Fund economics Ch. 3, 11, App. A, C | Fund size discipline Ch. 3, 14 | Fundraising funnel Ch. 15

G–L

Governance Ch. 19 | IRR App. A | IPO markets Ch. 4, 10, 21 | J-Curve App. B, C | LP due diligence Ch. 13 | LP pacing cycles Ch. 5, 12 | Liquidation preference Ch. 19, App. B

M–P

M&A exit strategy Ch. 4, 10 | Market cycle Ch. 21 | Moat stack Ch. 8 | MOIC Key Terms, App. A | NRR Ch. 9, App. A | Ownership discipline Ch. 11, 14 | Portfolio construction Ch. 12 | Power law returns Ch. 12, App. C | Pro-rata rights Ch. 12, 18, App. B

R–Z

Repricing (2022) Intro, Ch. 1 | Reserve strategy Ch. 12 | SAFE App. B | Series A gap Ch. 16 | Switching cost Ch. 7, 8 | Track record attribution Ch. 13 | TVPI Ch. 5, App. A | ValueAddVC.com Ch. 2 | Vertical AI Ch. 7, 8, 9 | Vintage year App. B | Waterfall App. B

About the Author

Trace Cohen was named after the Multiflow Trace, a VLIW minisupercomputer that pioneered parallel processing in the late 1980s. The company shipped about 120 machines and folded. The technology it invented now underpins modern computing—from GPU architectures to the cloud infrastructure powering today's AI. This tells you almost everything you need to know about his childhood.

His parents built and sold a large technology PR agency in New York, running communications for Sony and IBM at a time when those accounts represented the center of gravity of the tech industry. Trace grew up on computers, testing products, appearing in commercials, learning the difference between a beta and a launch build before most kids his age knew what a computer was.

He started his first company at seventeen, in his dorm room in 2007. He and his roommate built websites with Dreamweaver, HTML, and CSS, helping people get jobs online when SEO was still nascent and you needed a .edu email to join Facebook. That company ran for fourteen years before selling in 2021.

Along the way, he became member 400 of WeWork, joined New York Angels, and built his network during the formative years of the New York tech scene. He majored in entrepreneurship and technology at university, one of the early programs to take startup formation seriously as an academic discipline.

In recent years, Trace went deeper into building with AI, not just funding it. He built ValueAddVC.com using Claude in the terminal, connected to Vercel, Supabase, and Google Cloud, module by module, learning by doing. The platform gives founders, investors, and aspiring VCs the tools and frameworks usually locked behind years of experience or expensive relationships.

He takes meetings. He helps the next generation get further faster. He also learns from them, because in this community, the flow of insight goes in both directions when you show up with sharp elbows and an open mind.

He is still building.

Legal Notice

Copyright © 2026 Trace Cohen

All rights reserved. No part of this publication may be reproduced, distributed, or transmitted in any form without prior written permission of the author.

The information in this book is provided for educational purposes only. Nothing herein constitutes investment advice, and past performance is not indicative of future results.

First Edition, 2026

Printed in the United States of America